

Introduction to High Performance Computing

Autumn, 2018

Lecture 5

Git/Bitbucket notes

- After modifying files in your local repo, remember to *add*, *commit*, and *push* your changes
- Before starting work which will modify files in your local repo, remember to sync your fork and *pull* any changes into your local repo.
- This will help you avoid merging problems which occur when your fork and local repo are both being updated independently
- There is a link to a nice, short online tutorial on using git in the supplementary reading section of the course webpage.

Python notes

Main differences between arrays and lists:

Lists are *flexible*: heterogeneous data, can grow or shrink, numerical calculations can be slow/cumbersome

Arrays: calculations are generally faster, but elements must be homogeneous, difficult to adjust size

- Lists can be grown/shrunk more efficiently than arrays
- `L.append()` for a list, `L`, is more efficient than `np.append(A)` for an array, `A`. The numpy command will make a new copy of the array, the list will grow *in place*

Python notes

Getting comfortable with python:

- Have command of *all* of the material in online lectures –use exercises for self-assessment (solutions are online)
- Understand structure and purpose of functions (lecture 3 slides)
- Choose an editor + terminal combination for developing code. Can be *spyder* (distributed w/ anaconda), *canopy*, or *atom* + *jupyter qtconsole*. Use python3.x (e.g. python3.6)
- Understand *mysqrt.py* and *brown.py* (provided in lecture 4 and 5 directories of course repo)
- Further help: list of supplementary material on course webpage, office hours

Today

- Improving Brownian motion code
- Basics of optimization
- Using *scipy.minimize* package in python

Code development

- Use built-in functions whenever possible (e.g. `sqrt` instead of `mysqrt.sqrt2`)
- When working with arrays, *vectorize* code whenever possible:

```
In [8]: def sinx(x):  
...:     y = np.zeros_like(x)  
...:     for i,xi in enumerate(x):  
...:         y[i] = sin(xi)  
...:     return y  
...:
```

```
In [9]: x = np.random.randn(10000)
```

```
In [10]: timeit y1=sin(x) #vectorized  
159 µs ± 989 ns per loop (mean ± std. dev. of 7 runs, 10000 loops each)
```

```
In [11]: timeit y2 = sinx(x)  
18.6 ms ± 140 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```


Optimization with Python

We will focus almost entirely on the `scipy.optimize` package

- **`optimize` contains all standard methods for unconstrained optimization**
- **Choices for constrained optimization are more limited**
 - **Other packages are available (e.g. `pyopt`)**


Getting started – finding trends in data

Consider the noisy data:



Getting started – finding trends in data

Consider the noisy data:



- How can we construct a linear fit?
i.e. find c_1 and c_2 such that
$$T = c_1 x + c_2$$
- Python provides several options:
 - `scipy.stats`
 - `numpy.linalg.lstsq`
 - `scipy.optimize.curve_fit`
 - `numpy.polyfit`
 - Pandas, statsmodel, **probably quite a few more!**
- As a start, let's use `polyfit`

Getting started – finding trends in data

x and T are simple arrays:

```
In [8]: x = np.linspace(0,1,100)
```

```
In [9]: T = 2*x + 3 + 0.25*np.random.randn(100)
```

Getting started – finding trends in data

x and T are simple arrays:

```
In [8]: x = np.linspace(0,1,100)
```

```
In [9]: T = 2*x + 3 + 0.25*np.random.randn(100)
```

And we use polyfit to fit a 1st order polynomial (i.e. a line) to the data:


```
In [10]: C = np.polyfit(x,T,1)
```

```
In [11]: print "C1=",C[0]," , C2=",C[1]  
C1= 1.99587384666 , C2= 2.9703841544
```

```
In [13]: Tf = C[0]*x + C[1]
```

So the constants are close to what we expect, let's now display the fit:

Getting started – finding trends in data



- **polyfit constructs a linear least-squares fit to a n th-order polynomial:**

$$f(x) \approx c_1x + c_2x^2 + c_3x^3 + \dots + c_nx^n$$

so, it is only the fitting parameters that must be linear

- **The functional form for least-squares fits can be arbitrary (polynomial, trigonometric, etc...)**

General linear least squares

- Let's say you have data T_1, T_2, \dots, T_n at the points x_1, x_2, \dots, x_n and expect the data to fit a function of the form:

$$T = c_a + c_b g_b(x) + c_c g_c(x) + \dots$$

- Here, g_b, g_c, \dots can be any “well-behaved” functions. For polyfit, they have to be of the form x^n (with n an integer)
- As needed for *linear* least squares, the parameters c_a, c_b, \dots appear in a linear form
- Substituting each of the data points into the functional form (for the case with three parameters), we expect:

$$\begin{bmatrix} 1 & g_b(x_1) & g_c(x_1) \\ 1 & g_b(x_2) & g_c(x_2) \\ \vdots & \vdots & \vdots \\ 1 & g_b(x_n) & g_c(x_n) \end{bmatrix} \begin{bmatrix} c_a \\ c_b \\ c_c \end{bmatrix} \approx \begin{bmatrix} T_1 \\ T_2 \\ \vdots \\ T_n \end{bmatrix}$$

- And we can use linear least-square to find the “best” c_a, c_b , and c_c .

General linear least squares

$$\begin{bmatrix} 1 & g_b(x_1) & g_c(x_1) \\ 1 & g_b(x_2) & g_c(x_2) \\ \vdots & \vdots & \vdots \\ 1 & g_b(x_n) & g_c(x_n) \end{bmatrix} \begin{bmatrix} c_a \\ c_b \\ c_c \end{bmatrix} \approx \begin{bmatrix} T_1 \\ T_2 \\ \vdots \\ T_n \end{bmatrix}$$

- `np.linalg.lstsq` **requires the $n \times 3$ matrix on the left and the $n \times 1$ column vector on the right as input**

General linear least squares

- What does linear least-squares have to do with optimization?
- Rewrite matrix “equation” on previous slide as: $M\mathbf{c} \approx \mathbf{t}$
- **Least-squares problem is:** Find \mathbf{c} such that $\epsilon = (M\mathbf{c} - \mathbf{t})^2$ is *minimized*
- How do we minimize this?

General linear least squares

- What does linear least-squares have to do with optimization?
- Rewrite matrix “equation” on previous slide as: $M\mathbf{c} \approx \mathbf{t}$
- **Least-squares problem is:** Find \mathbf{c} such that $\epsilon = (M\mathbf{c} - \mathbf{t})^2$ is *minimized*
- How do we minimize this?
- **Single-variable calculus:** set derivative with respect to \mathbf{c} to zero, and check that the second derivative at the minimum is positive
- **Generalizing to n-dimensions:**
$$\begin{aligned}\epsilon &= (M\mathbf{c} - \mathbf{t})^T (M\mathbf{c} - \mathbf{t}) \\ &= \mathbf{c}^T M^T M \mathbf{c} - 2\mathbf{c}^T M^T \mathbf{t} + \mathbf{t}^T \mathbf{t} \\ \frac{\partial \epsilon}{\partial \mathbf{c}^T} &= 2M^T M \mathbf{c} - 2M^T \mathbf{t}\end{aligned}$$

General linear least squares

- What does linear least-squares have to do with optimization?
- Rewrite matrix “equation” on previous slide as: $M\mathbf{c} \approx \mathbf{t}$
- **Least-squares problem is:** Find \mathbf{c} such that $\epsilon = (M\mathbf{c} - \mathbf{t})^2$ is *minimized*
- How do we minimize this?
- **Single-variable calculus:** set derivative with respect to \mathbf{c} to zero, and check that the second derivative at the minimum is positive
- **Generalizing to n-dimensions:**
$$\begin{aligned}\epsilon &= (M\mathbf{c} - \mathbf{t})^T (M\mathbf{c} - \mathbf{t}) \\ &= \mathbf{c}^T M^T M \mathbf{c} - 2\mathbf{c}^T M^T \mathbf{t} + \mathbf{t}^T \mathbf{t} \\ \frac{\partial \epsilon}{\partial \mathbf{c}^T} &= 2M^T M \mathbf{c} - 2M^T \mathbf{t}\end{aligned}$$
- **Now setting the derivative to zero, \mathbf{c} must satisfy:** $M^T M \mathbf{c} = M^T \mathbf{t}$
- **All linear least-squares routines are solving this system of equations.**
 - **Routines for nonlinear least squares are also available**

Optimization: basic ideas

- For linear least squares, we don't need to check the second derivative (due to the general “shape” of norms like ϵ).
- For more general functions it is (often) useful to look at the second derivative
- Consider “general” optimization problem:

Find \mathbf{x} so that $f(\mathbf{x})$ is minimized

- To start, consider behavior near minimizing point, x^*
 - In 1-D: $f(x^* + h) = f(x^*) + \frac{df}{dx}|_{x^*}h + \frac{d^2f}{dx^2}|_{x^*}\frac{h^2}{2} + O(h^3)$

and since the derivative is zero at a minimum:

$$f(x^* + h) = f(x^*) + \frac{d^2f}{dx^2}|_{x^*}\frac{h^2}{2} + O(h^3)$$

So, the second derivative must be positive at a minimum (otherwise f is smaller at x^*+h than it is at x^*)

Optimization: basic ideas

- To start, consider behavior near minimizing point, \mathbf{x}^*

- **In 1-D:** $f(x^* + h) = f(x^*) + \frac{df}{dx}\bigg|_{x^*} h + \frac{d^2f}{dx^2}\bigg|_{x^*} \frac{h^2}{2} + O(h^3)$

- **Now in n-dimensions:**

$$f(x_i^* + h_i) = f(x_i^*) + \sum_{j=1}^n \frac{\partial f}{\partial x_j}\bigg|_{x_i^*} h_j + \sum_{j=1}^n \sum_{k=1}^n \frac{\partial^2 f}{\partial x_j \partial x_k}\bigg|_{x_i^*} \frac{h_j h_k}{2} + O(|h|^3)$$

or in vector notation:

$$f(\mathbf{x}^* + \mathbf{h}) = f(\mathbf{x}^*) + \mathbf{h}^T \nabla f|_{\mathbf{x}^*} + \frac{1}{2} \mathbf{h}^T H|_{\mathbf{x}^*} \mathbf{h} + O(|h|^3)$$

Here, H is the *Hessian*, $H_{j,k} = \frac{\partial^2 f}{\partial x_j \partial x_k}$, a $n \times n$ symmetric matrix.

- **At the minimum, the *gradient*, is zero, so we have:**

$$f(\mathbf{x}^* + \mathbf{h}) = f(\mathbf{x}^*) + \frac{1}{2} \mathbf{h}^T H|_{\mathbf{x}^*} \mathbf{h} + O(|h|^3)$$

- **The Hessian must be positive for any (small) \mathbf{h}**

Optimization: basic ideas

- All optimization routines require the user to specify the function to be minimized (usually called the *cost* function or the *objective* function)
- Most optimization routines use the gradient to search for a minimum (starting from an user-specified “guess”)
- Some routines use the Hessian, some don't.

Optimization: basic ideas

- All optimization routines require the user to specify the function to be minimized (usually called the *cost* function or the *objective* function)
- Most optimization routines use the gradient to search for a minimum (starting from an user-specified “guess”)
- Some routines use the Hessian, some don't
- Today, we'll look at two methods:
 1. (Truncated) Newton's method: uses gradient and Hessian
 2. BFGS method: uses gradient

Newton's method

- **Classical method:**

1. **Evaluate function, gradient and Hessian at a guess, \mathbf{x}_0**

2. **Use these values to fit a quadratic to the function:**

$$g(\mathbf{h}) = f(\mathbf{x}_0) + \mathbf{h}^T \nabla f|_{\mathbf{x}_0} + \frac{1}{2} \mathbf{h}^T H|_{\mathbf{x}_0} \mathbf{h}$$

3. **Setting $\frac{\partial g}{\partial \mathbf{h}} = 0$:**

$$H|_{\mathbf{x}_0} \mathbf{h} = -\nabla f|_{\mathbf{x}_0}$$

4. **This can be solved for \mathbf{h} and the new guess for the minimum is then:**

$$\mathbf{x}_1 = \mathbf{x}_0 + \mathbf{h}$$

- **Classical method only works well if guess is good**
- **In practice, a quadratic fit may be inaccurate or inappropriate**
 - **Truncated-Newton and BFGS methods both address this**

Truncated Newton's method

- **Classical method:**

1. Evaluate function, gradient and Hessian at a guess, x_0
2. Use these values to fit a quadratic to the function:

$$g(\mathbf{h}) = f(\mathbf{x}_0) + \mathbf{h}^T \nabla f|_{\mathbf{x}_0} + \frac{1}{2} \mathbf{h}^T H|_{\mathbf{x}_0} \mathbf{h}$$

3. Setting $\frac{\partial g}{\partial \mathbf{h}} = 0$:

$$H|_{\mathbf{x}_0} \mathbf{h} = -\nabla f|_{\mathbf{x}_0}$$

4. This can be solved for \mathbf{h} and the new guess for the minimum is then:
 $\mathbf{x}_1 = \mathbf{x}_0 + \mathbf{h}$

Truncated approach:

1. Iteratively solve for \mathbf{h} as in the classical approach (but truncate iterations if local curvature is negative)
2. Then search in the direction of \mathbf{h} for the minimum along this direction
 - Step 2 is called a *line search*, idea is similar to finding the zero of a 1-d function (e.g. Newton-Rhapson, secant, Brent methods)

Optimizers in Scipy

- **We will use the `scipy.optimize` package**
- **And particularly, the function, `scipy.optimize.minimize`**
- **The optimization *method* can be specified when calling this function, e.g.:**
`method = 'Newton-CG'`
- **This will use truncated Newton's method**
- **The CG indicates the conjugate gradient method is used to solve for h**
- **There is a separate CG method for optimization**


Optimizers in Scipy

- **Newton-CG requires the user to specify:**
 - **A python function which computes the cost function**
 - **A function for computing the gradient**
 - **A function for computing the Hessian (or instruct the optimizer to approximate it)**
 - **A guess, x_0 , where the optimizer starts its search for a minimum**

Simple example

- Simple illustrative example:

Find x and y that minimize $f = -\exp(-\alpha(x - x_0)^2 - \beta(y - y_0)^2)$



Simple example

The code, *gauss2d.py*, applies a few different approaches to this problem

- **It has three functions:** `gauss2d`, `gauss2d_grad`, `gauss2d_hess`
- **`gauss2d_grad` returns the two components of the gradient:**

$$\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) = (-2\alpha(x - x_0)f, -2\beta(y - y_0)f)$$

- **`gauss2d_hess` returns the 2 x 2 Hessian matrix:**

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix} = \begin{bmatrix} -2\alpha(f + (x - x_0)\frac{\partial f}{\partial x}) & 4\alpha\beta(x - x_0)(y - y_0)f \\ 4\alpha\beta(x - x_0)(y - y_0)f & -2\beta(f + (y - y_0)\frac{\partial f}{\partial y}) \end{bmatrix}$$

- **The minimizer is called near the bottom with different inputs**

Simple example

- The code, *gauss2d.py*, applies a few different approaches
- Let's first look at the Newton-CG, approximate gradient case (`ncg1=True`)
- The call to the minimizer is:

```
minimize(gauss2d,xguess,args=parameters,method='Newton-CG',jac=gauss2d_grad)
```

Note: since the name of the Hessian function is not specified, it is approximated from the gradient

Simple example

- **Running the code:**

```
In [44]: run gauss2d
Gauss2d, Newton-CG, approximate Hessian
optimum: [ 0.5    0.25]
info:      status: 0
  success: True
    njev: 17
    nfev: 6
    fun: -1.0
      x: array([ 0.5 ,  0.25])
message: 'Optimization terminated successfully.'
  nhev: 0
    jac: array([ -5.37315725e-09,  -2.68657863e-09])
```

Notes:

- **gauss2d sets (x_0, y_0) to $(0.5, 0.25)$ so the answer is correct**
- **17 gradient evaluations and 6 function evaluations were needed**
- **What happens if we now specify the exact Hessian (`ncg2=True`)?**

Simple example

Running the code with:

```
minimize(gauss2d,xguess,args=parameters,method='NewtonCG',  
        jac=gauss2d_grad,hess=gauss2d_hess)
```

```
Gauss2d, Newton-CG, exact Hessian  
optimum: [ 0.5   0.25]  
info:      status: 0  
  success: True  
    njev: 9  
    nfev: 6  
    fun: -1.0  
      x: array([ 0.5 ,  0.25])  
message: 'Optimization terminated successfully.'  
    nhev: 4  
    jac: array([-2.28865316e-09, -1.14432658e-09])
```

Notes:

- The number of gradient evaluations reduced from 17 to 9
- gauss2d also has options for BFGS and Nelder-Mead methods – how do these work?

Overview of BFGS

- BFGS is generally the method-of-choice for unconstrained optimization problems involving smooth functions
- Works surprisingly well for non-smooth/noisy functions as well
- Can be memory-intensive, use L-BFGS-B for large problems
- Motivating idea:
 - Computing and inverting the Hessian can be very expensive
 - Instead, approximate the (inverse) Hessian, and update this approximation at each step (quasi-Newtonian methods)

Overview of BFGS

- **Motivating idea:**
 - Computing and inverting the Hessian can be very expensive
 - Instead, approximate the (inverse) Hessian, and update this approximation at each step (quasi-Newtonian methods)
- **Newton-CG:**
 1. **Solve for h:** $H|_{\mathbf{x}_0} \mathbf{h} = -\nabla f|_{\mathbf{x}_0}$
 2. **Search along direction of h for minimum $\rightarrow \mathbf{x}_1$ and repeat step 1**
- **BFGS:**
 1. **Approximate inverse Hessian and solve for h:** $M^i \approx H^{-1}, \mathbf{h} = -M^i \nabla f|_{\mathbf{x}_i}$
 2. **Similar to Newton-CG**
 3. **Update M:** $M^{i+1} = f(\tilde{M}^i, h)$
- **The update should result in a positive-definite (approximate) Hessian, and the BFGS update formula produces particularly good results**
- **Now back to *gauss2d***

Simple example -- BFGS

Results for BFGS with approximate gradient

```
minimize(gauss2d,xguess,args=parameters,method='BFGS',jac=False)
```

```
Gauss2d, BFGS, approximate gradient
optimum: [ 0.49999556  0.24999778]
info:      status: 0
    success: True
        njev: 5
        nfev: 20
hess_inv: array([[ 0.60077592, -0.19961208],
                 [-0.19961208,  0.90019394]])
    fun: -0.9999999999753861
        x: array([ 0.49999556,  0.24999778])
message: 'Optimization terminated successfully.'
    jac: array([ -8.85874033e-06, -4.43309546e-06])
```

Simple example -- BFGS

Results for BFGS with exact gradient

```
minimize(gauss2d,xguess,args=parameters,method='BFGS',jac=gauss2d_grad)
```

```
Gauss2d, BFGS, exact gradient
optimum: [ 0.49999557  0.24999779]
info:      status: 0
    success: True
        njev: 5
        nfev: 5
hess_inv: array([[ 0.60077589, -0.19961205],
                 [-0.19961205,  0.90019397]])
    fun: -0.9999999999754784
        x: array([ 0.49999557,  0.24999779])
message: 'Optimization terminated successfully.'
    jac: array([ -8.85828506e-06,  -4.42914253e-06])
```

Notes:

1. **BFGS** requires less function, gradient evaluations than Newton-CG
2. For large (slow) problems, specifying the gradient (and Hessian if appropriate) can make a big difference