

Introduction to High Performance Scientific Computing

Autumn, 2018

Lecture 9

Avoiding Fortran

- **Libraries**
- **Introduction to F2Py**

Using libraries with Fortran

- **Avoid writing own code whenever possible!**
- **Many well-established libraries for scientific computing are freely available.**
 - **Save time: don't have to write/test own code**
 - **Standard libraries have been extensively tested**
 - **Libraries often optimized to run fast, difficult to do better**

Using libraries with Fortran

- Avoid writing own code whenever possible!
- Many well-established libraries for scientific computing are freely available.
 - Save time: don't have to write/test own code
 - Standard libraries have been extensively tested
 - Libraries often optimized to run fast, difficult to do better

Examples:

- Netlib: minpack, odepack, blas, ...
- FFTW – fastest Fourier transform in west
- Lapack (we will focus on this)

Usually possible to call libraries written in *c* from *fortran* and vice versa

Examples

minpack: ...software for solving nonlinear equations and nonlinear least squares problems, netlib.org/minpack (fortran 77)

blas: The BLAS (Basic Linear Algebra Subprograms) are routines that provide standard building blocks for performing basic vector and matrix operations. netlib.org/blas (fortran 77, see also atlas, <http://math-atlas.sourceforge.net/> for optimized blas)

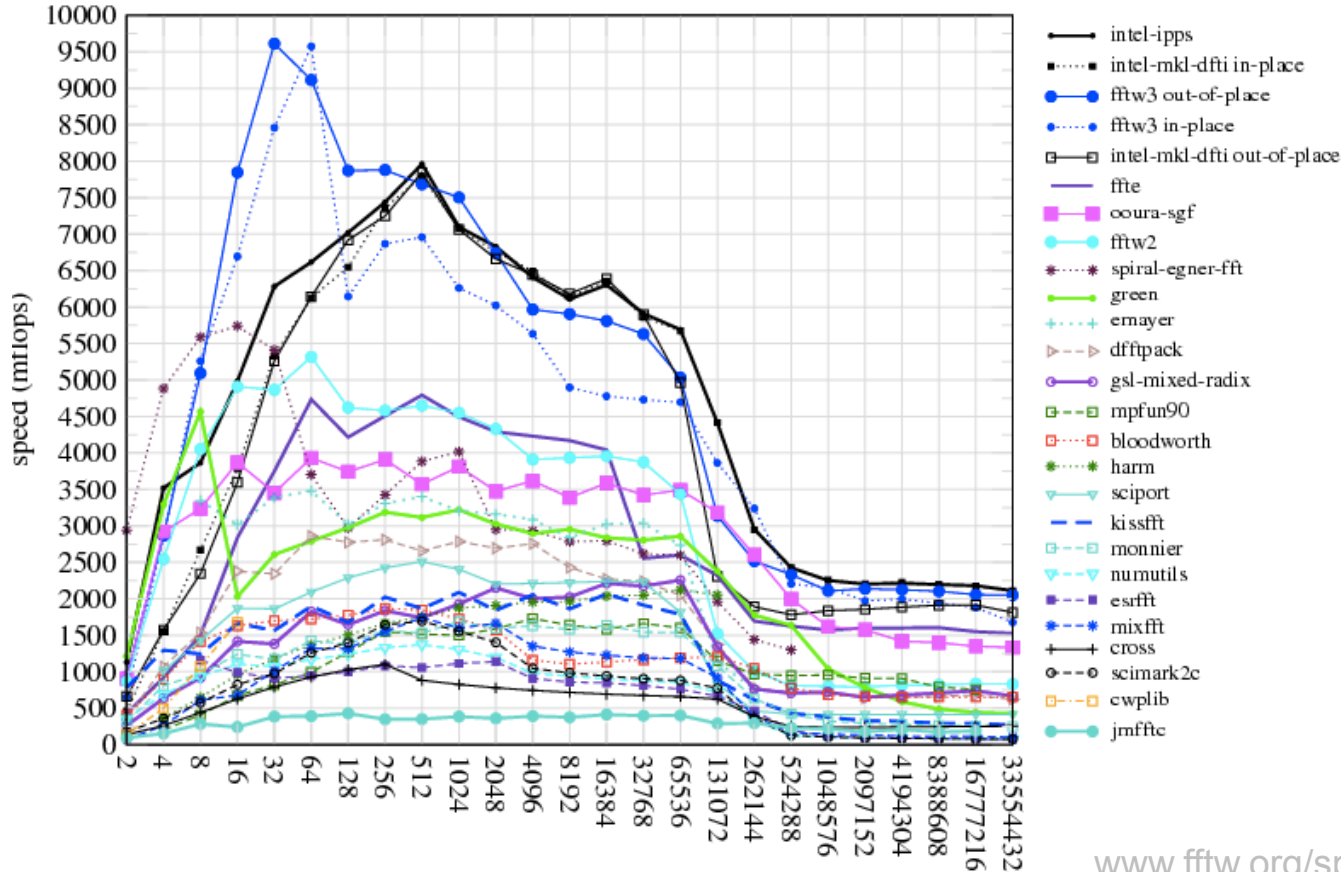
fftw: fastest Fourier transform in the west, FFTW is a free collection of fast C routines for computing the Discrete Fourier Transform in one or more dimensions. fftw.org

Choosing the right library

- Sometimes there are many libraries for the same task
- Example: “Right” fast fourier transform package depends on:
compiler, architecture, programmer’s background

double-precision complex, 1d transforms

powers of two



Lapack

netlib.org/lapack: LAPACK provides routines for solving systems of simultaneous linear equations, least-squares solutions of linear systems of equations, eigenvalue problems, and singular value problems.

- **You've probably used lapack without realizing it!**
 - **It is built into Numpy, Matlab, Maple, and R**
 - **It (and blas) are included with mac os x**
 - **It is written in Fortran 77, and using it takes some getting used to**

Using lapack

1. Find the driver subroutine you want to use (browse lapack site or use google)
 - E.g. DGESV to solve linear system of equations
 - Understanding lapack naming convention helps:

All driver and computational routines have names of the form XYYZZZ, where for some driver routines the 6th character is blank.

The first letter, X, indicates the data type as follows:

S	REAL
D	DOUBLE PRECISION
C	COMPLEX
Z	COMPLEX*16 or DOUBLE COMPLEX

from lapack user guide

Using lapack

1. Find the driver subroutine you want to use (browse lapack site or use google)
 - E.g. DGESV to solve linear system of equations
 - Understanding lapack naming convention helps:

All driver and computational routines have names of the form XYYZZZ, where for some driver routines the 6th character is blank.

The next two letters, YY, indicate the type of matrix

BD bidiagonal

DI diagonal

GB general band

GE general (i.e., unsymmetric, in some cases rectangular)

GG general matrices, generalized problem (i.e., a pair of general matrices)

GT general tridiagonal

HB (complex) Hermitian band

HE (complex) Hermitian

Note: this is just an excerpt, there are 28 possible matrix types in total.

Using lapack

1. Find the driver subroutine you want to use (browse lapack site or use google)
 - E.g. DGESV to solve linear system of equations
 - Understanding lapack naming convention helps:

All driver and computational routines have names of the form XYYZZZ, where for some driver routines the 6th character is blank.

The next two letters, YY, indicate the type of matrix

BD bidiagonal
DI diagonal
GB general band
GE general (i.e., unsymmetric, in some cases rectangular)
GG general matrices, generalized problem (i.e., a pair of general matrices)
GT general tridiagonal
HB (complex) Hermitian band
HE (complex) Hermitian

Note: this is just an excerpt, there are 28 possible matrix types in total.

So, DGESV is a double precision routine for SolVing (SV) systems with general matrices

Using lapack

2. Look at documentation for subroutine to see what input is needed and how output is returned

- googling “lapack dgesv” takes me to:

http://www.netlib.org/lapack/explore-html/d8/d72/dgesv_8f.html

- Eight variables in subroutine header (are they input and/or output?)

Function/Subroutine Documentation

```
subroutine dgesv ( integer                                N,  
                  integer                                NRHS,  
                  double precision, dimension( lda, * ) A,  
                  integer                                LDA,  
                  integer, dimension( * )                IPIV,  
                  double precision, dimension( ldb, * ) B,  
                  integer                                LDB,  
                  integer                                INFO  
                  )
```

DGESV computes the solution to system of linear equations $A * X = B$ for GE matrices

Using lapack

2. Look at documentation for subroutine to see what input is needed and how output is returned

- googling “lapack dgesv” takes me to:
http://www.netlib.org/lapack/explore-html/d8/d72/dgesv_8f.html
- Eight variables in subroutine header (are they input and/or output?)

Parameters

[in]	N	<pre>N is INTEGER The number of linear equations, i.e., the order of the matrix A. N >= 0.</pre>
[in]	NRHS	<pre>NRHS is INTEGER The number of right hand sides, i.e., the number of columns of the matrix B. NRHS >= 0.</pre>
[in,out]	A	<pre>A is DOUBLE PRECISION array, dimension (LDA,N) On entry, the N-by-N coefficient matrix A. On exit, the factors L and U from the factorization A = P*L*U; the unit diagonal elements of L are not stored.</pre>

- N and NRHS are input, A is input which is modified so the LU decomposition is returned.

Using lapack

2. Look at documentation for subroutine to see what input is needed and how output is returned

- googling “lapack dgesv” takes me to:
http://www.netlib.org/lapack/explore-html/d8/d72/dgesv_8f.html
- Eight variables in subroutine header (are they input and/or output?)
- Solving $A X = B$:
 - A must be a LDA x N matrix (though only the 1st N rows will be used)
 - B is LDB x NRHS (again only 1st N rows are used)
 - X will be N x NRHS and must be extracted from B
 - IPIV contains information about the LU decomposition

Using lapack

3. Create test example to check routine is used correctly:

dgesv_test.f90: !example illustrating use of lapack routine, dgesv, to solve $AX=B$

1. Create (nonsingular) matrix A and RHS, B
2. Send copies of A and B to *dgesv* (why use copies?)
3. Extract the solution x
4. Verify that $Ax = B$

Using lapack

3. Create test example to check routine is used correctly:

dgesv_test.f90: !example illustrating use of lapack routine, dgesv, to solve $AX=B$

1. Create (nonsingular) matrix A and RHS, B

```
!allocate arrays, create matrices A and B
allocate(A(LDA,N),B(LDB,NRHS),IPIV(N))
allocate(Atemp(size(A,1),size(A,2)),Btemp(LDB,NRHS))
A(1,:) = (/1.0,2.0,3.0,4.0/)
A(2,:) = (/4.0,3.0,2.0,1.0/)
A(3,:) = A(1,):**2
A(4,:) = sqrt(A(2,:))

B(:,1) = (/ -2.0,2.0,-1.0,1.0/)
```

Using lapack

3. Create test example to check routine is used correctly:

dgesv_test.f90: !example illustrating use of lapack routine, dgesv, to solve $AX=B$

2. Send copies of A and B to *dgesv* (why use copies?)

```
!solve Ax = B
  Atemp = A
  Btemp = B
  call dgesv(N, NRHS, Atemp, LDA, IPIV, Btemp, LDB, INFO)
  print *, 'INFO=', INFO
```


Using lapack

3. Create test example to check routine is used correctly:

dgesv_test.f90: !example illustrating use of lapack routine, dgesv, to solve $AX=B$

3. Extract the solution x

```
!extract soln from Btemp  
  allocate(x(N,NRHS))  
  x = Btemp(1:N,:)
```

Using lapack

3. Create test example to check routine is used correctly:

dgesv_test.f90: !example illustrating use of lapack routine, dgesv, to solve $AX=B$

4. Verify that $Ax = B$

```
print *, 'test:', matmul(A,x)-B
```

Compiling and linking to libraries

Generally, if *program.f90* uses library *libname*:

```
$ gfortran -c program.f90  
$ gfortran -o program.exe program.o -llibname
```

or:

```
$ gfortran -o program.exe program.f90 -llibname
```

Note: compiler will look for *libname* (often a file, *liblibname.a*) in “standard” locations (e.g. */usr/lib*)

If you have built the library elsewhere, specify the location with “**-L *libpath***” option when linking. Here *libpath* specifies location of library (e.g. */users/prasun/somewhere*)

Compiling and linking to libraries

So, for dgesv:

```
$ gfortran -o dgesv_test.exe dgesv_test.f90 -llapack  
$ ./dgesv_test.exe
```

```
INFO=          0  
test: 2.1316282072803006E-014  1.7763568394002505E-015  2.8421709430404007E-014  
-5.3290705182007514E-015
```

What does INFO=0 indicate?

Intro to F2Py

- **Major projects:**
 - **Variety of tasks, some more computationally intensive than others**
 - **F2Py:**
 - **Place expensive parts in Fortran routines (why?)**
 - **Everything else in Python**
 - **Use F2Py to convert Fortran routines into Python modules which are called from main Python code**

F2Py example

- Simple example: call fortran function *sumxy* from python

```
!-----  
!function to sum two numbers provided as input  
!note that the function name is a variable whose  
!type should be declared in the header  
function sumxy(x,y)  
    implicit none  
    real(kind=8), intent(in) :: x,y  
    real(kind=8) :: sumxy  
  
    sumxy = x + y  
end function sumxy
```

```
!-----
```

F2Py example

Simple example: call fortran function *sumxy* from python (see f2pyfunction1.f90)

```
!-----  
!function to sum two numbers provided as input  
!note that the function name is a variable whose  
!type should be declared in the header  
function sumxy(x,y)  
    implicit none  
    real(kind=8), intent(in) :: x,y  
    real(kind=8) :: sumxy  
  
    sumxy = x + y  
end function sumxy  
!-----
```

Use f2py to “compile” this function:

```
$ f2py -c f2pyfunction1.f90 -m f1
```

```
$ ls f1.so  
f1.so
```

f1.so is a *shared object file* which can be imported in python

F2Py example

Import *f1.so* in python:

```
In [119]: import f1
```

```
In [120]: f1?
```

```
Type:      module
```

```
String form: <module 'f1' from 'f1.so'>
```

```
File:      ~/Desktop/fortran_dev/f1.so
```

```
Docstring:
```

```
This module 'f1' is auto-generated with f2py (version:2).
```

```
Functions:
```

```
    sumxy = sumxy(x,y)
```


F2Py example

Import *f1.so* in python:

```
In [119]: import f1
```

```
In [120]: f1?
```

```
Type:      module
```

```
String form: <module 'f1' from 'f1.so'>
```

```
File:      ~/Desktop/fortran_dev/f1.so
```

```
Docstring:
```

```
This module 'f1' is auto-generated with f2py (version:2).
```

```
Functions:
```

```
    sumxy = sumxy(x,y)
```

```
In [121]: f1.sumxy(3,5)
```

```
Out[121]: 8.0
```

F2Py example

Fortran function, *sumxy*, now available as python function, *f1.sumxy*

```
In [122]: f1.sumxy?  
Type:      fortran  
String form: <fortran sumxy>  
Docstring:  
sumxy = sumxy(x,y)
```

Wrapper for ``sumxy``.

Parameters

x : input float

y : input float


Returns

sumxy : float

F2Py example

- Subroutines work similarly
- Important to use *intent(out)* to specify what python function will return

```
!-----  
!subroutine to sum two numbers provided as input  
subroutine sumxy2(x,y,sumxy)  
  implicit none  
  real(kind=8), intent(in) :: x,y  
  real(kind=8), intent(out) :: sumxy  
  
  sumxy = x + y  
  
end subroutine sumxy2
```



```
$ f2py -c f2pyfunction1.f90 -m f2
```

F2Py example

- Subroutines work similarly
- Important to use *intent(out)* to specify what python function will return

```
$ f2py -c f2pyfunction1.f90 -m f2
```

```
In [128]: import f2
```

```
In [129]: f2.sumxy2?
```

```
Type:          fortran
```

```
String form: <fortran object>
```

```
Docstring:
```

```
sumxy = sumxy2(x,y)
```

```
Wrapper for ``sumxy2``.
```

```
Parameters
```

```
-----
```

```
x : input float
```

```
y : input float
```

```
Returns
```

```
-----
```

```
sumxy : float
```

F2Py and other wrappers

- Fortran has built-in capability for calling c routines
- F2Py can be used to call c routines from python
- But cython is a more commonly-used c-python interface
- See also swig.org for interfaces between c/C++ and other languages

Fortran timing functions

- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:

```
!timing variables
real(kind=8) :: cpu_t1,cpu_t2,clock_time
integer(kind=8) :: clock_t1,clock_t2,clock_rate
```

Fortran timing functions

- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:

```
!timing variables
  real(kind=8) :: cpu_t1,cpu_t2,clock_time
  integer(kind=8) :: clock_t1,clock_t2,clock_rate

  call system_clock(clock_t1)
  call cpu_time(cpu_t1)
  !loop over intervals computing each interval's contribution to
  integral
... Midpoint quadrature ...
  call cpu_time(cpu_t2)
  print *, 'elapsed cpu time (seconds) =',cpu_t2-cpu_t1

  call system_clock(clock_t2,clock_rate)
  print *, 'elapsed wall time (seconds)= ',
                                dble(clock_t2-clock_t1)/dble(clock_rate)
```

Fortran timing functions

- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:

```
$ ./midpoint_t.exe
elapsed cpu time (seconds) = 8.6239999999999997E-003
elapsed wall time (seconds)= 9.12799966E-03
N= 512000
sum= 3.1415926535901515
error= 3.5837999234900053E-013
```

- Can place timing commands throughout code to find bottlenecks

Fortran timing functions

- Also, often have a *theoretical estimate* of how cost scales with problem size
- A method may require $O(N)$ (or $O(N \ln_2 N)$ or $O(N^2)$) operations
- But does your implementation of the algorithm match theory?
- How do compiler optimizations affect performance?
- Carefully timing code while varying the problem size can help answer these questions