

Introduction to High Performance Scientific Computing

Autumn, 2018

Lecture 10

Announcements

- **HW 2 will be posted around 8pm today**
- **HW1 feedback will be provided around Wednesday next week (solutions will be posted at the end of this week)**
- **A new VM is being installed on computer in Huxley 408 and 410. Should be substantial improvement on MLC VMs**
 - **Will post announcement when they are ready (hopefully sometime tomorrow!)**

Today

- **A few comments on Fortran**
- **F2Py with fortran modules**
- **Basic computer architecture**

Notes on Fortran

- Do not attempt to guess Fortran syntax! – look it up instead
- Do not develop Fortran code in Python using f2py – get a Fortran-only code working first
- Test code (compile and run) after adding a few (4-5) lines of code. Don't write 100 lines and then test!
- Lecture 7: code structure, variable types, loops, if-then, subroutines
- Lecture 8: allocatable arrays, functions modules
- Lecture 9: Lapack, f2py
 - Note: On VMs, you may need to use “f2py3” instead of “f2py”

Notes on Fortran

- Fortran is case *insensitive*, the variables `a` and `A` are treated as the same variable
- Python is case sensitive, `a` and `A` are treated as different variables
- This can cause problems when using `f2py`. Safe approach: use all lowercase letters for variable names
- Slicing/indexing are a little different in Fortran and Python
 - Fortran: counting starts from 1, the first `n` elements of `A` are `A(1:n)`.

Notes on Fortran

Debugging code

- If the compiler is giving you a series of errors, look at the topmost error message. It will include a line number
- If the code runs but crashes with a segmentation fault: probably a problem with array indices (e.g. trying to access the 12th element of a dimension(11) array)
- If you can't tell where the code is crashing, add print statements, e.g. `print *, 1` --some code-- `print *, 2`
When the code is run, if the 1 prints to screen but not the 2, you know where the problem is.
- If the code runs, but gives the wrong answer, add print statements outputting values of variables, try a small problem size where you know what values the variables can take

Notes on Fortran

- Modules consist of:
 1. Module variables
 2. Module sub-programs
 - Module variables are “available” throughout the module
 - They *do not* need to be declared within sub-programs
 - They *do not* need to be provided as input/output in the sub-program header
 - Module variables and module sub-programs are also “available” in any program or sub-program that *uses* the module
 - A module by itself doesn’t do anything
 - There should be a “main” program which uses it
 - You can compile a module by itself: `gfortran -c module.f90`
 - But to generate an executable, you need a program:
`gfortran -o program.exe module.f90 program.f90`
- When re-compiling, first remove the previous .mod file

Notes on Fortran

- **Compiling code that uses lapack routines:**

```
$ gfortran -c program.f90  
$ gfortran -o program.exe program.o -llapack  
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

Similarly, with f2py:

```
$ f2py -llapack -c program.f90 -m module_name
```


F2Py and fortran modules

- F2Py will recognize subroutines and functions in modules
- What about variables?
 - Try f2py with circle module from last week (*f2pymodule_circle.f90*)

```
$ f2py -c f2pymodule_circle.f90 -m cmod
```

```
In [10]: import cmod
```

```
In [11]: cmod.<tab>  
cmod.circle  cmod.so
```

Need to look at *cmod.circle*

F2Py and fortran modules

Need to look at *cmod.circle*:

```
In [12]: cmod.circle?
```

Docstring:

```
'd'-scalar  
initialize_pi()
```

```
Wrapper for ``initialize_pi``.
```

```
circumference = circumference(radius)
```

```
Wrapper for ``circumference``.
```

Parameters

```
radius : input float
```

Returns

```
circumference : float  
area = area(radius)
```

And similar info for “area”

F2Py and fortran modules

How do we access variables and methods in *cmod.circle*?

```
In [23]: cmod.circle.<tab>
cmod.circle.area          cmod.circle.circumference
cmod.circle.initialize_pi  cmod.circle.pi
```

Can initialize *pi* in python:

```
In [9]: cmod.circle.pi
Out[9]: array(0.0)

In [10]: cmod.circle.pi = pi

In [11]: cmod.circle.pi
Out[11]: array(3.141592653589793)
```

F2Py and fortran modules

How do we access variables and methods in *cmmod.circle*?

```
In [23]: cmmod.circle.<tab>
cmmod.circle.area          cmmod.circle.circumference
cmmod.circle.initialize_pi  cmmod.circle.pi
```

Can initialize *pi* in python:

```
In [9]: cmmod.circle.pi
Out[9]: array(0.0)

In [10]: cmmod.circle.pi = pi

In [11]: cmmod.circle.pi
Out[11]: array(3.141592653589793)
```

Can also initialize allocatable arrays, see *f2pymodule_circle.f90...*

F2Py and fortran modules

Can also initialize allocatable arrays, see *f2py module_circle_array.f90*:

```
!module for computing circumference, area, and "mass" of circle
module circle
  implicit none
  real(kind=8) :: pi
  real(kind=8), allocatable, dimension(:) :: weights, mass
  save
```

...

...

```
subroutine compute_mass(radius, mass)
  !compute mass = weights*area
  implicit none
  real(kind=8), intent(in) :: radius
  real(kind=8), intent(out) :: mass(:)

  mass = weights*area(radius)
end subroutine compute_mass
```

F2Py and fortran modules

Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

F2Py and fortran modules

Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

F2Py and fortran modules

Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

3. Initialize variables:

```
In [15]: cmoda.circle.pi=pi
```

```
In [16]: cmoda.circle.weights=arange(5)
```

```
In [17]: cmoda.circle.pi,cmoda.circle.weights
```

```
Out[17]: (array(3.141592653589793), array([ 0.,  1.,  2.,  3.,  4.]))
```


F2Py and fortran modules

Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

3. Initialize variables:

```
In [15]: cmoda.circle.pi=pi
```

```
In [16]: cmoda.circle.weights=arange(5)
```

```
In [17]: cmoda.circle.pi,cmoda.circle.weights
```

```
Out[17]: (array(3.141592653589793), array([ 0.,  1.,  2.,  3.,  4.]))
```

4. Compute mass:

```
In [18]: cmoda.circle.compute_mass(2.0)
```

```
In [19]: cmoda.circle.mass
```

```
Out[19]: array([ 0. , 12.56637061, 25.13274123, 37.69911184,
50.26548246])
```

F2Py and fortran modules

Exercise: How does f2py treat the subroutine below with weights now provided as input? Can n be removed from the input?

```
subroutine compute_mass(n,weights,radius)
  !compute mass = weights*area
  implicit none
  integer :: n
  real(kind=8), dimension(n), intent(in) :: weights
  real(kind=8), intent(in) :: radius

  if (.not. allocated(mass)) allocate(mass(size(weights)))
  mass = weights*area(radius)

end subroutine compute_mass
```

Notes on compiling

- Up to now:

```
$ gfortran -c program.f90  
$ gfortran -o program.exe program.o -llapack  
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

- But typically want to turn on *optimization* -O flag:

```
$ gfortran -O3 -c program.f90
```

- -O3 is highest level of optimization (can also use -O1, -O2)

Notes on compiling

- Look at `midpoint_time.f90` compiled with and without `-O3`:

```
$ ./mt2.exe
elapsed cpu time (seconds) = 0.32510499999999998
elapsed wall time (seconds)= 0.326231003
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010

$ ./mt2_03.exe
elapsed cpu time (seconds) = 0.14376099999999997
elapsed wall time (seconds)= 0.144617006
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010
```

- Optimization can make a substantial difference
- `f2py` uses `-O3` by default

Vectorization

- Many built-in functions in Fortran can operate on arrays
 - e.g. with an array (or matrix), x , can compute $\exp(x)$ all at once
- We know it is important to avoid loops in Python, does the same apply to Fortran?

Vectorization

- Many built-in functions in Fortran can operate on arrays
 - e.g. with an array (or matrix), x , can compute $\exp(x)$ all at once
- We know it is important to avoid loops in Python, does the same apply to Fortran?
- Consider a simple example: computing the sine of a large array, $\sin(x)$

- **Python:**

```
#compute sin(x) using loop
def sin_loop(x):
    s = np.zeros_like(x)
    for i,xi in enumerate(x):
        s[i] = np.sin(xi)
    return s
```

```
#vectorized
def sin_vec(x):
    s = np.sin(x)
    return s
```

Vectorization

- Many built-in functions in Fortran can operate on arrays
 - e.g. with an array (or matrix), x , can compute $\exp(x)$ all at once
- We know it is important to avoid loops in Python, does the same apply to Fortran?
- Consider a simple example: computing the sine of a large array, $\sin(x)$

- **Python:**
sine_array.py

```
#compute sin(x) using loop
def sin_loop(x):
    s = np.zeros_like(x)
    for i,xi in enumerate(x):
        s[i] = np.sin(xi)
    return s
```

```
#vectorized
def sin_vec(x):
    s = np.sin(x)
    return s
```

Vectorization

- **Python:** `In [12]: x = np.random.randn(500000)`
`In [13]: timeit sin_loop(x)`
808 ms \pm 32.2 ms per loop (mean \pm std. dev. of 7 runs, 1 loop each)

`In [15]: timeit sin_vec(x)`
11.2 ms \pm 346 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)
- **Vectorized version is 70x faster**
- **The increase in speed does depend on the array size**
- **What about Fortran?**

Vectorization

- What about Fortran? See `sine_array.f90`
- Compile module with `f2py`: `$ f2py -c sine_array.f90 -m s1`

Import and run in python:

```
In [25]: from s1 import sine_array as sa
```

```
In [26]: sa.x = x
```

```
In [27]: n = x.size
```

```
In [28]: timeit sa.sin_loop(n)  
10.5 ms  $\pm$  245  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100  
loops each)
```

```
In [29]: timeit sa.sin_vec(n)  
10.6 ms  $\pm$  192  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100  
loops each)
```

Vectorization

- What about Fortran? See `sine_array.f90`
 - Vectorization is *essential* in Python (or Matlab)
 - Much less important in Fortran
 - Provided you optimize properly when compiling!