

Introduction to High Performance Scientific Computing

Autumn, 2018

Lecture 12

Announcements

- **Tutorials from this week on will be in Huxley 410**
 - Except last Tuesday of last week of term
 - Will give quick demo using the new VMs at beginning of this week's lab
 - See class webpage for instructions on launching these VMs
- **Tuesday office hours also in Huxley 410 (except last week of term)**
- A correction was added to the homework assignment on Sunday (related to definition of testing error)
- Homework 1 marking is in-progress, hope to provide scores+feedback by Thursday afternoon
- If you haven't done so already, please fill out the short course feedback form here: <https://goo.gl/forms/q0Vq81pu1tbZCNlf1>

Comments

- **Fortran and f2py tips/notes**
 - If you're running a Fortran routine from qtconsole, the output from the Fortran code will print to the *Unix* terminal where qtconsole is launched
 - If using ipython (or ipython3) terminal, the Fortran output will be displayed within the ipython terminal
 - A segmentation fault is usually associated with arrays and array sizes, e.g. trying to access the 11th element in a 10-element array
 - Compiling with the -fcheck=bounds flag may give a more helpful error message:
 - gfortran -fcheck=bounds -o test.exe test.f90
 - f2py --opt='-fcheck=bounds' -c test.f90 -m t1
- **Be careful copying and pasting code from slides! Sometimes powerpoint changes text in strange ways!**

Today

Introduction to parallel computing (carried over from lecture 11)

Introduction to OpenMP

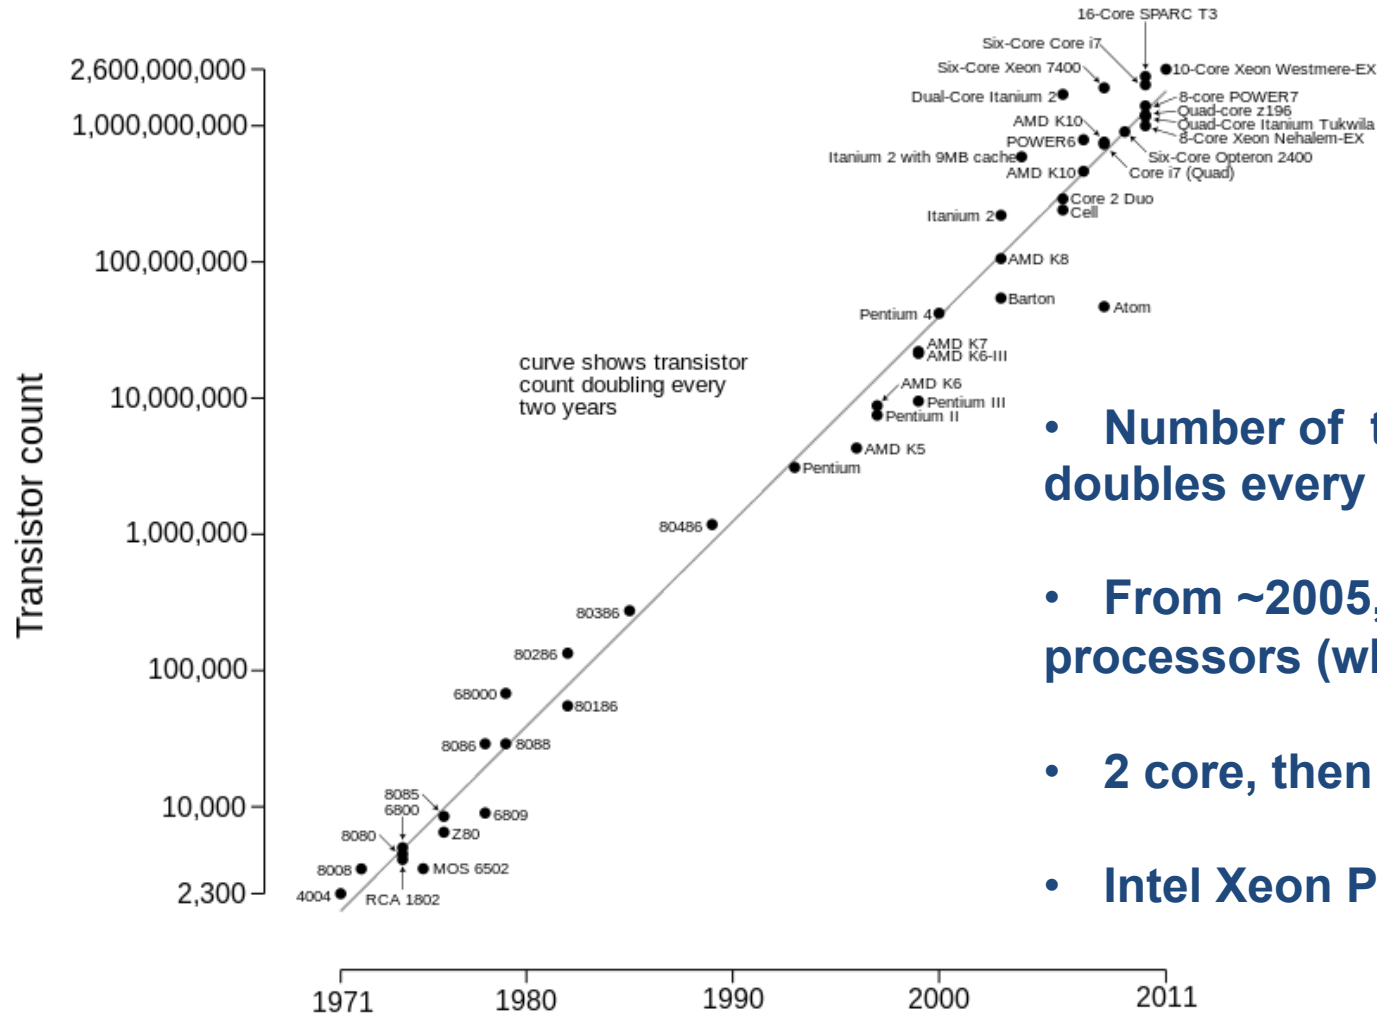
Getting started

Parallel regions

Parallel loops

Moore's law

Microprocessor Transistor Counts 1971-2011 & Moore's Law




- Number of transistors on chip doubles every two years
- From ~2005, multicore processors (why?)
- 2 core, then 4, 6, 8, 16
- Intel Xeon Phi: 60+ cores!

Algorithms and hardware


Speed-Up
Factor

Derived from Computational Methods



Speed-Up
Factor

Derived from Supercomputer Hardware




Why parallelize a code?

1. Serial (single-processor) code is too slow

or

2 Serial code is too big


Parallel computing paradigms



Shared memory

- One 4-core chip with shared memory (RAM)
- MPI can coordinate communication between cores
- OpenMP generally easier to use for shared-memory systems
- MPI = *Message Passing Interface*
- OpenMP = *Open Multi-Processing*


Parallel computing paradigms



Distributed memory

- Each (4-core) chip has its own memory
- The chips are connected by network ‘cables’
- MPI coordinates communication between two or more CPUs


Parallel computing paradigms



Related approaches:

- Hybrid programming: mix of shared-memory (OpenMP) and distributed-memory (MPI) programming
- GPU's: Shared memory programming (CUDA or OpenCL)
- Coprocessors and co-array programming

Example: computing an integral



- Estimate integral with midpoint rule,

$$I = \int_{x_0}^{x_6} f(x) dx$$

1. Compute:

$$f(x_1^*), f(x_2^*), \dots$$


2. Compute areas of rectangles:

$$I_1 = (x_1 - x_0) * f(x_1^*)$$

3. Sum areas:


$$I \approx I_1 + I_2 + I_3 + \dots$$

Example: computing an integral




- How to parallelize?
- With three processors, can compute areas of two rectangles on each processor
- Not practical for small calculations, but could split $1e7$ rectangles across, say, 10 processors

Scaling and performance



- How do we measure performance of a parallel code?
- Serial code: Optimize the *efficiency* → cost required to obtain a certain level of accuracy

Scaling and performance



- How do we measure performance of a parallel code?
- Serial code: Optimize the *efficiency* → cost required to obtain a certain level of accuracy
- Parallel code: Also optimize *scaling* or *speedup*: how much faster is the calculation when the number of procs is increased?

Speedup

- **Speedup = Computation time on one proc/time on N procs = T_s/T_p**
- **Ideal: N = 10 processors, speedup = N = 10**

Speedup

- **Speedup = Computation time on one proc/time on N procs = T_s/T_p**
- **Ideal: $N = 10$ processors, speedup = $N = 10$**
- **Real life: Speedup will be less than N (possibly much less) Why?**
 - **Startup costs**
 - **Communication**
 - **Only part of the algorithm parallelizes**
- **Typically interested in performance of large problems running on large number of processors**
 - **Workstation: $N = 16, 32$**
 - **Imperial HPC (cx2): $N = 256+$**
 - **UK HPC (Archer): $N = 1e3, 1e4, \dots$**
- **Ahmdal's law provides guidance**

Ahmdal's law

- Usually only part of a computation can be parallelized
 - One processor: $T(1) = s + p$
 - Two processors: $T(2) = s + p/2$
 - N processors: $T(N) = s + p/N$

p is the part of the code that can be parallelized

Ahmdal's law

- Usually only part of a computation can be parallelized
 - One processor: $T(1) = s + p$
 - Two processors: $T(2) = s + p/2$
 - N processors: $T(N) = s + p/N$

p is the part of the code that can be parallelized

So, if only half the code can be parallelized ($s = p = 0.5$), Then the maximum speedup $T(1)/T(N \rightarrow \infty) = (s+p)/(s) = 2$

It is important for p to be much larger than s !

Ahmdal's law

Speedup $T(1)/T(N) = (s+p)/(s+p/N)$

Example: $s = 0.1$, $p = 0.9$

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

Ahmdal's law

Speedup $T(1)/T(N) = (s+p)/(s+p/N)$

Example: $s = 0.1$, $p = 0.9$

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

Waste of resources to use $N=256$!

Strong and weak scaling

- **Strong scaling:** Time needed to solve a problem of fixed size as number of processors increases
- **Weak scaling:** Time needed for problem with *fixed size per processor*

Overview

- OpenMP provides a fairly easy approach to parallelizing c/c++ or fortran code
- Add *directives* indicating how/where the code should run in parallel
- Parallel regions have multiple threads, each of which should be assigned computational tasks
 - OpenMP is for *shared-memory* parallel programming
 - Each thread has access to all variables that existed before parallel region was created
 - This can cause problems if multiple threads try to change the same variable!
- Particularly useful for parallelizing loops
- When compiling, add *-fopenmp* flag

Overview

Program starts with
single *master thread*

Start program



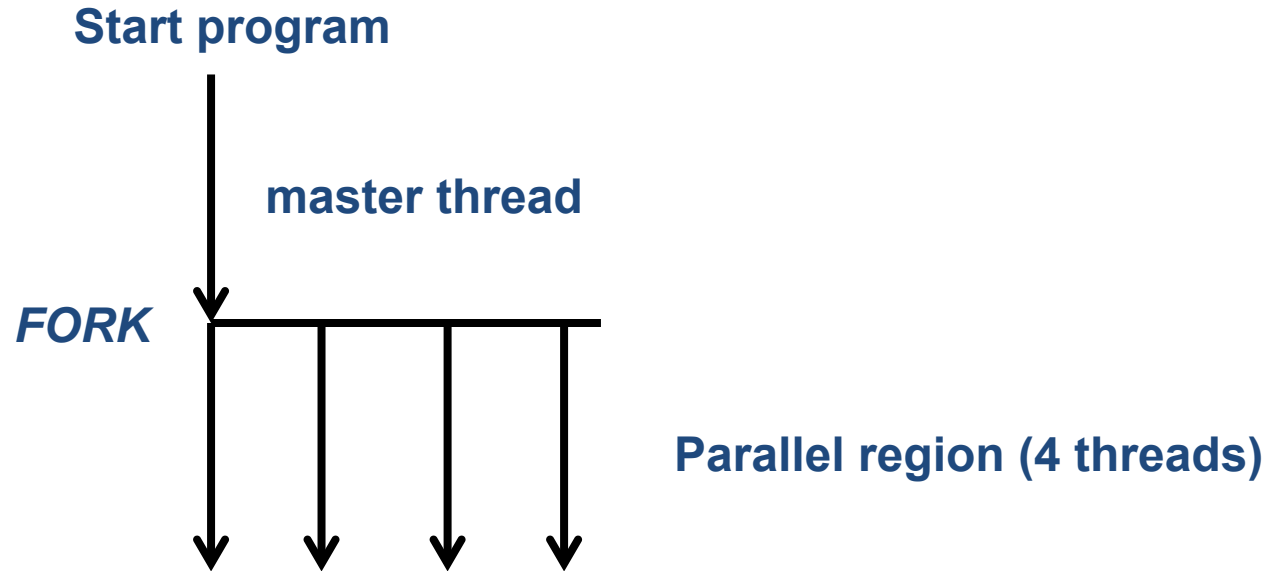
master thread

Overview

Program starts with
single *master thread*

Then, launch parallel
region with multiple
threads.

Each thread has
access to all
variables introduced
previously



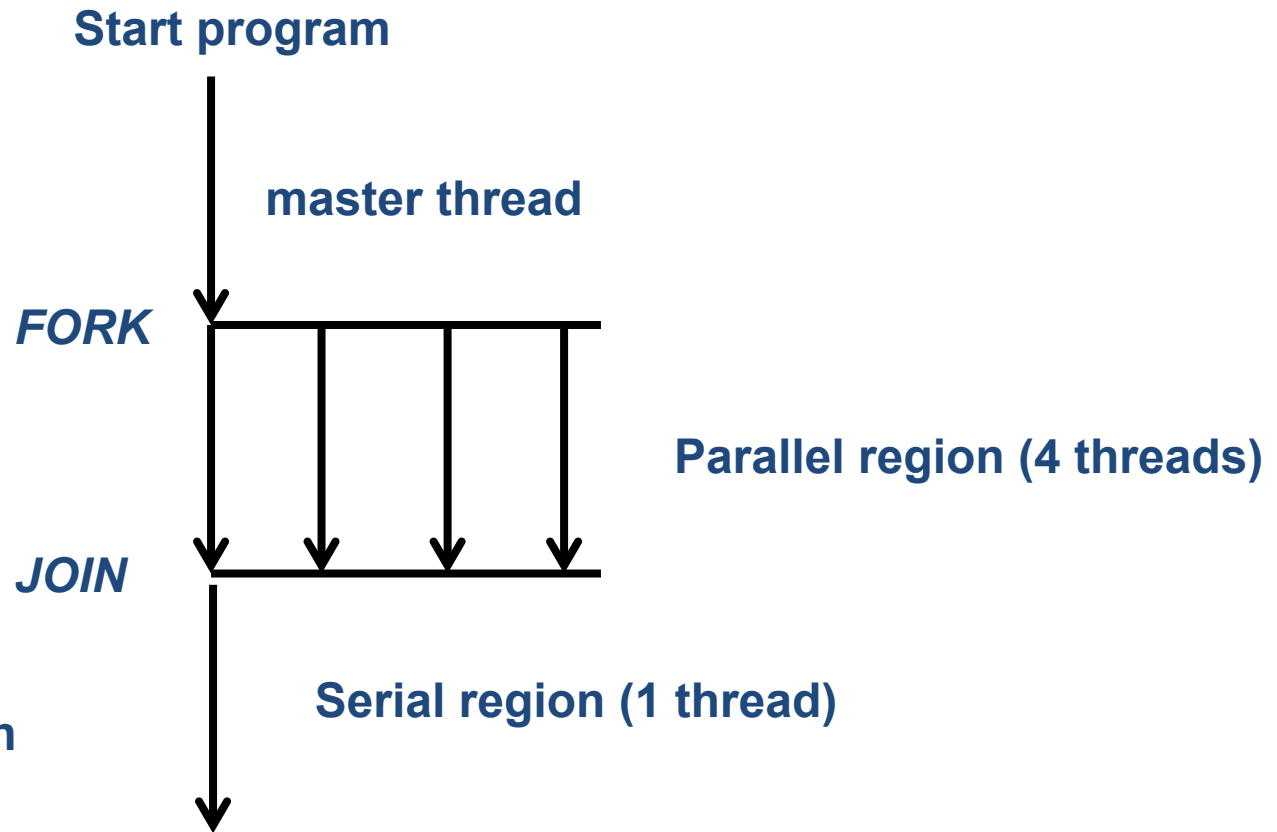
Overview

Program starts with single *master thread*

Then, launch parallel region with multiple threads.

Each thread has access to all variables introduced previously

Can end parallel region if/when desired and launch parallel regions again in future as needed



Simple OpenMP example

- Launch parallel region, get info on threads (see *firstomp_v0.f90*)
- Must use openMP module, *omp_lib*
 - This makes functions like *omp_get_num_threads* available

!Getting started with OpenMP

```
program firstomp
  use omp_lib !makes OpenMP routines, variables available
  implicit none
  integer :: NumThreads, threadID
```

Simple OpenMP example

- Launch parallel region, get info on threads (see *firstomp_v0.f90*)
- Must use openMP module, *omp_lib*
 - This makes functions like *omp_get_num_threads* available

!Getting started with OpenMP

```
program firstomp
  use omp_lib !makes OpenMP routines, variables available
  implicit none
  integer :: NumThreads, threadID
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  print *, 'this is thread', threadID, ' of ', NumThreads
!$OMP END PARALLEL
```

- !\$OMP starts an OpenMP directive (#pragma omp in c)
- !\$OMP PARALLEL starts a parallel region (forks a number of threads)
 - *omp_get_num_threads* tells us how many threads are forked
 - *omp_get_thread_num* tells us which thread is being used

Simple OpenMP example

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**

Simple OpenMP example

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**
- **Remember: *threadID* is a *shared* variable.**
 - **Each thread is writing to the same variable with it's id, so only the “last” thread has its ID displayed**

Simple OpenMP example

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**
- **Remember: *threadID* is a *shared* variable.**
 - **Each thread is writing to the same variable with it's id, so only the “last” thread has its ID displayed**
- **How can we fix this? First approach: define a *critical* region...**

Simple OpenMP example

- A critical region, defined with `!$OMP CRITICAL`, runs in serial
 - The threads carry out their tasks sequentially (*firstomp_v1.f90*)

```
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  !$OMP CRITICAL
    threadID = omp_get_thread_num()
    print *, 'this is thread', threadID, ' of ', NumThreads
  !$OMP END CRITICAL
!$OMP END PARALLEL
```

Simple OpenMP example

- A critical region, defined with `!$OMP CRITICAL`, runs in serial
 - The threads carry out their tasks sequentially (*firstomp_v1.f90*)

```
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  !$OMP CRITICAL
    threadID = omp_get_thread_num()
    print *, 'this is thread', threadID, ' of ', NumThreads
  !$OMP END CRITICAL
!$OMP END PARALLEL
```

So now, if we compile and run:

```
$ ./testv1.exe
this is thread          0  of          4
this is thread          2  of          4
this is thread          1  of          4
this is thread          3  of          4
```


Simple OpenMP example

- Critical regions useful for: data I/O, displaying results
- But forcing serial execution in a parallel region, not generally desirable
- Better approach: set `threadID` to be a *private* variable (*firstomp.f90*)
 - Then, each thread will have their own private copy of the variable

Simple OpenMP example

- Critical regions useful for: data I/O, displaying results
- But forcing serial execution in a parallel region, not generally desirable
- Better approach: set threadID to be a *private* variable (*firstomp.f90*)
 - Then, each thread will have their own private copy of the variable

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  print *, 'this is thread', threadID, ' of ', NumThreads
!$OMP END PARALLEL
```

```
$ ./test.exe
this is thread      0  of      4
this is thread      2  of      4
this is thread      1  of      4
this is thread      3  of      4
```

Simple parallel calculation

- Can use *threadID* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
```

```
  if (threadID==0) then
    call subroutine1(in1,out1)
  elseif (threadID==1) then
    call subroutine1(in2,out2)
  end if
```

```
!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)

Overview

- **OpenMP (primarily) consists of *directives* and *routines***
- **Directives are denoted with !OMP**
 - !OMP parallel, !OMP critical, ...
 - **Directives are recognized when `–fopenmp compile-flag` is used**
 - **Otherwise, they are interpreted as comments**
 - **What happens if you use:**
!\$ print *, “compiled with `–fopenmp`”
- **Routines are available via the `use omp_lib` command**
 - **e.g. `omp_get_thread_num` and `omp_get_num_threads`**

Parallel loops

- Loops form the backbone of most scientific codes
- They should be parallelized whenever possible
- They can be parallelized if the calculations of each iteration are independent of each other (no data dependencies)

```
do i1 = 1,n  
    x(i1) = y(i1) + z(i1)  
end do
```

Ok to parallelize

```
do i1 = 1,n  
    norm = norm + abs(x(i1))  
end do
```

Can't parallelize easily: each thread updating, *norm*

Parallel loops

- OpenMP makes it very easy to parallelize loops

```
!$OMP parallel do
do i1 = 1,n
    x(i1) = y(i1) + z(i1)
end do
!$OMP end parallel do
```

- OpenMP automatically distributes iterations across threads
 - If NumThreads=2 and n=10, iterations 1,...,5 would be given to thread 0 and iterations 6,...,10 would be done by thread 1 (or vice versa)
 - The iterated variable, *i1*, is automatically set to *private*. Each thread has its own copy.

Parallel loops

- Simple example (*loop_omp1.f90*):

```
!$OMP parallel do private(threadID)
do i1 = 1,size(x)
    x(i1) = y(i1) + z(i1)
    threadID = omp_get_thread_num()
    print *, 'iteration ',i1,' assigned to thread ',threadID
end do
!$OMP end parallel do

print *, 'test:', maxval(abs(x-y-z))
```

- Note: *threadID* again set to *private*
- Compile and run...

Parallel loops

- Simple example (*loop_omp1.f90*):

```
!$OMP parallel do private(threadID)
do i1 = 1,size(x)
    x(i1) = y(i1) + z(i1)
    threadID = omp_get_thread_num()
    print *, 'iteration ',i1,' assigned to thread ',threadID
end do
!$OMP end parallel do
```

```
print *, 'test:', maxval(abs(x-y-z))
```

```
$ gfortran -fopenmp -o testl1.exe loop_omp1.f90
```

```
$ ./testl1.exe
```

iteration	1	assigned to thread	0
iteration	3	assigned to thread	2
iteration	2	assigned to thread	1
iteration	4	assigned to thread	3
test:	2.2204460492503131E-016		