

High Performance Computing

Autumn, 2018

Lecture 13

Notes

- **Compiling:** the operating system where you compile and create and executable or “object” is the same operating system where you have to use the executable/object
 - Mac OS X, Windows, Xubuntu, and Oracle Linux are all *distinct* operating systems
- **Fortran:** avoid using integers in mathematical expressions
 - If a and b are integers:
 - **bad:** (a/b) – this results in an integer!
 - **good:** $(\text{dble}(a)/\text{dble}(b))$
 - **Similarly should use** $(2.\text{d}0/3.\text{d}0)$ **instead of** $2/3$

Today

More on OpenMP

More on loops

Reductions

Setting number of threads

A few useful OpenMP commands

Last time: parallel loops

- OpenMP makes it very easy to parallelize loops

```
!$OMP parallel do
do i1 = 1,n
    x(i1) = y(i1) + z(i1)
end do
!$OMP end parallel do
```

- OpenMP automatically distributes iterations across threads
 - If NumThreads=2 and n=10, iterations 1,...,5 would be given to thread 0 and iterations 6,...,10 would be done by thread 1 (or vice versa)
 - The iterated variable, *i1*, is automatically set to *private*. Each thread has its own copy.

Parallel loops: nested loops

- Should we parallelize the inner or outer loop? (assuming $M \sim N$)

```
!$OMP parallel do private(i1)
do j1 = 1,N
  do i1 = 1,M
    x(i1,j1) = y(i1,j1) + z(i1,j1)
  end do
end do
!$OMP end parallel do
```

Better to parallelize outer thread (setting the inner variable, $i1$, to private)

If inner loop is parallelized: forking/joining of threads is repeated with each outer loop: this is inefficient!

Last time: Parallel loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(i1)
do j1 = 2,N
    do i1 = 1,M
        x(i1,j1) = x(i1,j1-1)
    end do
end do
!$OMP end parallel do
```

Is the order of the iterations important? (data dependency)

Do different iterations assign values to same variable? (race condition)

Parallel loops: nested loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(j1)
do i1 = 1,M
    do j1 = 2,N
        x(i1,j1) = x(i1,j1-1)
    end do
end do
!$OMP end parallel do
```

Correct

- Solution: swap inner and outer loops
- Now, computation of x is “safe.”
 - The “ $i1$ loop” is parallelized, and calculations of x do not depend on the order in which $i1$ is iterated.

Example: computing a norm

Simple example: computing norm: $\text{sum}(|x|)$

Serial version:

```
do i1 = 1, size(x)
    norm = norm + abs(x(i1))
end do
```


Parallel loops

1. Let each thread have it's own copy of norm

```
norm=0.d0
partial_norm=0.d0
!$OMP parallel firstprivate(partial_norm),private(threadID)
!$OMP do
do i1 = 1,size(x)
    partial_norm = partial_norm + abs(x(i1))
end do
!$OMP end do
```

- The partial sum, *partial_norm*, is a private variable which must be initialized
- *firstprivate* initializes each thread's value to the value set before the parallel region

Example: computing a norm

Simple example: computing norm: $\sum(|x|)$

Serial version:

```
do i1 = 1, size(x)
    norm = norm + abs(x(i1))
end do
```

Parallel version:

```
!$OMP parallel firstprivate(partial_norm)
!$OMP do
do i1 = 1, size(x)
    partial_norm = partial_norm + abs(x(i1))
end do
!$OMP end do

!$OMP critical
norm = norm + partial_norm
!$OMP end critical
!$OMP end parallel
```

Example: computing a norm

- Typically want to avoid using critical regions
- *reduction* provides a simpler approach:

```
!$OMP parallel do reduction(+:norm)
do i1 = 1,size(x)
    norm = norm + abs(x(i1))
end do
!$OMP end parallel do
```

Example: computing a norm

- Typically want to avoid using critical regions
- *reduction* provides a simpler approach (*omp_norm2.f90*):

```
!$OMP parallel do reduction(+:norm)
do i1 = 1,size(x)
    norm = norm + abs(x(i1))
end do
!$OMP end parallel do
```

- Generally, *reduction* “reduces” an array of numbers distributed across multiple threads to a single number
- Several operations are available, a few common operators are: +,-,*,max,min,.and,.or.
- Not specific to OpenMP! In MPI, we will use *MPI_REDUCE*.
- Due to ease-of-use and usefulness, one of the most important tools in parallel computing!

Example: reduction with *min*

- Here, computation of x is parallelized
- Reduction is used to find $\min(|x|)$

```
!$OMP parallel do reduction(min:xmin)
do i1=1,size(x)
    x(i1) = z(i1)+y(i1)
    xmin = min(abs(x(i1)),xmin)
end do
!$OMP end parallel do
```

Setting number of threads

- By default, OpenMP “detects” the number of threads on computer and uses all of them

- Can also set threads in two ways:

1. Within code with *omp_set_num_threads*, e.g.:

```
!$ call omp_set_num_threads(2)
```

(the “!\$” ensures this is only called if *-fopenmp* flag is used when compiling)

2. From Unix terminal before program execution:

```
$ export OMP_NUM_THREADS=2
```

Other useful OpenMP directives

- Consider a parallel region of code:

```
!$OMP parallel
```

```
!code run by *each* thread
```

```
!$OMP end parallel
```

- There are a number of directives which we can use in the parallel region

Other useful OpenMP directives

- **do-loops**

```
!$OMP parallel private(i1)

do i1=1,N
    !some operations
end do

!$OMP end parallel
```

- In the example above, the full do-loop is run by each thread

Other useful OpenMP directives

- **do-loops**

```
!$OMP parallel private(i1)

do i1=1,N
    !some operations
end do

!$OMP end parallel
```

- **In the example above, the full do-loop is run by each thread**

```
!$OMP parallel private(i1)
!$OMP do
do i1=1,N
    !some operations
end do
!$OMP end do
!$OMP end parallel
```

- **Now, this is the same as a parallel do loop**

Other useful OpenMP directives

- **Sections**

```
!$OMP parallel
```

```
!$OMP sections
```

```
!$OMP section  
    !code run by one thread
```

```
!$OMP section  
    !code run by second thread
```

```
!$OMP section  
    !code run by another thread
```

```
!$OMP end sections
```

```
!$OMP end parallel
```

- **Manually assign tasks to threads**
- **For example, invert four matrices (of the same size)**
- **Could have four “sections”, one for each matrix inversion**

Simple parallel calculation

- Can use *threadID* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()

  if (threadID==0) then
    call subroutine1(in1,out1)
  elseif (threadID==1) then
    call subroutine1(in2,out2)
  end if

!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)

Simple parallel calculation

- Can use *sections* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  !$OMP sections
    !$OMP section
      call subroutine1(in1,out1)
    !$OMP section
      call subroutine1(in2,out2)
  !$OMP end sections
!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)

Other useful OpenMP directives

- **Single**

```
!$OMP parallel
```

```
!$OMP single
```

```
    !code run by only one thread
```

```
!$OMP end single
```

```
!$OMP end parallel
```

- **Used to run commands only once within parallel region**
- **Useful for: print statements, data input/output**

Other useful OpenMP directives

- **Single**

```
!$OMP parallel
```

```
!$OMP single
```


```
    !code run by only one thread
```

```
!$OMP end single nowait ←
```

```
!$OMP end parallel
```


- Used to run commands only once within parallel region
- Useful for: print statements, data input/output
- Add *nowait* tag to allow other threads to continue while one thread is in *single* region

Example: computing an integral



- How to parallelize?
- With three processors, can compute areas of two rectangles on each processor
- Not practical for small calculations, but could split $1e7$ rectangles across, say, 10 processors

Example: computing an integral



- How to parallelize?
- With three processors, can compute areas of two rectangles on each processor
- Not practical for small calculations, but could split $1e7$ rectangles across, say, 10 processors
- This is a simple *reduction* problem

Example: computing an integral

- Serial version:

```
!loop over intervals computing each interval's contribution to integral
do i1 = 1,N
    xm = dx*(dble(i1)-0.5d0) !midpoint of interval i1
    call integrand(xm,f)
    sum_i = dx*f
    sum = sum + sum_i !add contribution from interval to total
end do
```

- Parallel version (see *midpoint_omp.f90*):

```
!$OMP parallel do private(xm,f,sum_i),reduction(+:sum)
do i1 = 1,N
    xm = dx*(dble(i1)-0.5d0) !midpoint of interval i1
    call integrand(xm,f)
    sum_i = dx*f
    sum = sum + sum_i !add contribution from interval to total
end do
!$OMP end parallel do
```

Example: computing an integral

- Is there any actual performance gain?
 - Use `system_clock` and `omp_set_num_threads` (see *midpoint_time_omp.f90*)

- $N=1000$

`numThreads = 1` wall time= 2.300000005E-04

`numThreads = 2` wall time= 6.97000010E-04

`numThreads = 4` wall time= 1.09699997E-03

- Here, parallelization slows down the calculation! Why?
- Recall Amdahl's law, here $s > p$
- s/p will change as N increases...

Example: computing an integral

- Is there any actual performance gain?
 - Use `system_clock` and `omp_set_num_threads` (see *midpoint_time_omp.f90*)

- $N=1e7$

`numThreads = 1` wall time= 1.11312997

`numThreads = 2` wall time= 0.6055430174

`numThreads = 4` wall time= 0.565499008

- Now, we see improved performance
- Speedup from two threads = 1.8
- No meaningful gain from four threads – laptop only has two cores

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- *Barriers* keep the threads synchronized:

```
!$OMP parallel
```

```
!Some code
```

```
!$OMP barrier
```

```
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- **Barriers** keep the threads synchronized:

```
!$OMP parallel
```

```
!Some code
```

```
!$OMP barrier
```

```
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier
- There are *implicit* barriers at end of !\$OMP do and !\$OMP single blocks

Thread-safe routines

- What happens when you call sub-program from within parallel region?
- Each thread will call it's own “copy” of sub-program
 - All “local” variables declared within sub-program are private to thread

```
!$OMP parallel
call sub1(in1,in2,out1,out2)
!$OMP end parallel
!-----
subroutine sub1(in1,in2,out1,out2)
  use mod1
  implicit none
  real(kind=8) intent(in) :: in1,in2
  real(kind=8) intent(out) :: out1,out2
  real(kind=8) :: local1

  !should not modify mod1 variables
  !out1,out2 should (usually) be
  !private in the calling parallel region

end subroutine sub1
```

Basic questions:

1. Does code give same answer independent of the total number of threads?
2. Is it independent of the *order* in which threads call the subroutine

If yes, the subroutine is *thread-safe*

Should not include OMP directives in subroutine called from within parallel region