

High Performance Computing

Autumn, 2018

Lecture 14

OpenMP

What directives/routines have been covered?

OpenMP

Directives:

!\$OMP parallel

!\$OMP parallel do

!\$OMP do

!\$OMP critical

!\$OMP single

!\$OMP sections

Routines:

omp_get_num_threads

omp_set_num_threads

omp_get_thread_num

reduction, private, firstprivate

Notes

Homeworks:

- **HW1:** Will respond to e-mail queries on marking by end of term (probably the last week of term)
 - Scores may go up or down
- **HW2:** Solutions should be posted end of this week, marks end of next week
- **HW3:** Posted today around 8pm
 - Due next ****Thursday**** (so you have extra time for final project)
 - Will require f2py + fortran + openmp (Try today's lecture code!)
 - Will need ffmpeg to save animations
 - Installation instructions will be posted on course webpage
 - Demo during this week's lab on Huxley 410 VM

XUbuntu VMs:

- Installed directly on Huxley 408 and 410 machines
- Available via software hub in MLC

Today

A little more on OpenMP: synchronization, thread-safe subroutines, nested for loops

Programming example: from PDE → algorithm → serial code → parallel code

Parallel loops: nested loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(j1)
do i1 = 1,M
    do j1 = 2,N
        x(i1,j1) = x(i1,j1-1)
    end do
end do
!$OMP end parallel do
```

Correct

- Solution: swap inner and outer loops
- Now, computation of x is “safe.”
 - The “ $i1$ loop” is parallelized, and calculations of x do not depend on the order in which $i1$ is iterated.

Parallel loops: nested loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do collapse(2)
do i1 = 1,M
    do j1 = 2,N
        x(i1,j1) = sin(y(i1,j1))
    end do
end do
!$OMP end parallel do
```

- Nested loops can be “collapsed”
 - If both loops are parallelizable
 - And there is no code “between” the loops

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- *Barriers* keep the threads synchronized:

```
!$OMP parallel
```

```
!Some code
```

```
!$OMP barrier
```

```
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- **Barriers** keep the threads synchronized:

```
!$OMP parallel
```

```
!Some code
```

```
!$OMP barrier
```

```
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier
- There are *implicit* barriers at end of !\$OMP do and !\$OMP single blocks

Thread-safe routines

- What happens when you call sub-program from within parallel region?
- Each thread will call it's own “copy” of sub-program
 - All “local” variables declared within sub-program are private to thread

```
!$OMP parallel
call sub1(in1,in2,out1,out2)
!$OMP end parallel
!-----
subroutine sub1(in1,in2,out1,out2)
  use mod1
  implicit none
  real(kind=8) intent(in) :: in1,in2
  real(kind=8) intent(out) :: out1,out2
  real(kind=8) :: local1

  !should not modify mod1 variables
  !out1,out2 should (usually) be
  !private in the calling parallel region

end subroutine sub1
```

Basic questions:

1. Does code give same answer independent of the total number of threads?
2. Is it independent of the *order* in which threads call the subroutine

If yes, the subroutine is *thread-safe*

Should not include OMP directives in subroutine called from within parallel region

Vectorizing code

In general (Python, Fortran, Matlab,...), avoid for loops and *vectorize* calculations involving arrays.

Example:

```
In [27]: x=np.linspace(0,1,101)
```

```
In [28]: f = np.empty_like(x)
```

```
In [29]: for i in range(size(x)):
.....:         f[i] = cos(x[i])**2
.....:
```

```
In [30]: f = cos(x)**2    Vectorized version of loop
```

Vectorizing code

In general (Python, Fortran, Matlab,...), avoid for loops and *vectorize* calculations involving arrays.

Example:

```
In [27]: x=np.linspace(0,1,101)
```

```
In [28]: f = np.empty_like(x)
```

```
In [29]: for i in range(size(x)):
.....:         f[i] = cos(x[i])**2
.....:
```

```
In [30]: f = cos(x)**2      Vectorized version of loop
```

- **Vectorized code will usually be faster, sometimes *much faster* in interpreted languages**
- **Exception: parallelizing Fortran code with OpenMp:**

vectorized code → loops → parallel loops

Programming example

Task: Compute temperature distribution in a room

Programming example

Task: Compute temperature distribution in a room

Governing equation: Heat equation (diffusion equation):

$$\frac{\partial T}{\partial t} = \alpha \nabla^2 T + S(\mathbf{x}, t)$$

$$T(\mathbf{x}, t = 0) = f(\mathbf{x}) \quad \text{Initial condition}$$

Here, S is a *heat source*. Boundary conditions should also be specified as appropriate.

Problem: given the source, initial condition, and boundary conditions, solve for the temperature distribution, $T(\mathbf{x}, t)$

Programming example

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

Initial condition

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

Boundary conditions

$$0 \leq x \leq 1$$

Programming example

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

Initial condition

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

Boundary conditions

$$0 \leq x \leq 1$$

First consider steady problem, e.g., $S = S(x)$, a and b are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x, t) = 0$$

Poisson equation

Programming example

First consider steady problem, e.g., $S = S(x)$, a and b are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0$$

Notes:

- 1. This is an extremely simple problem, easy to write down the analytical solution**
- 2. No need to use compiled language**
- 3. Certainly no need to parallelize**
- 4. But what about two-dimensional or three-dimensional problems?**
 - Then, the picture changes considerably!**
- 5. We are just considering the 1-D problem for illustrative purposes**

Programming example

First consider steady problem, e.g., $S = S(x)$, a and b are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 1, 2, \dots, N$$

$$(N + 1) * \Delta x = 1$$

Programming example

First consider steady problem, e.g., $S = S(x)$, a and b are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

With boundary conditions: $T_0 = T_a, \quad T_N = T_b$

Programming example

Equation for T_i :
$$\frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} = -S_i$$

In matrix form: $AT = b$

$$A = \begin{bmatrix} -2 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -2 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -2 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \vdots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -2 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -2 \end{bmatrix}, b = \Delta x^2 \begin{bmatrix} -T_a/\Delta x^2 - S_1 \\ -S_2 \\ \vdots \\ -S_i \\ \vdots \\ -S_{N-1} \\ -T_b\Delta x^2 - S_N \end{bmatrix}$$

- In 1-D, this is just a tridiagonal system of equations
- Easy to solve directly (with, say, *DGTSV*)

Programming example

- In two or three dimensions, A loses its simple banded structure
- Then, direct solution becomes very expensive for large N
- *Iterative* methods are a popular alternative

Programming example

- In two or three dimensions, A loses its simple banded structure
- Then, direct solution becomes very expensive for large N
- *Iterative* methods are a popular alternative
- Basic idea: rewrite $Ax=b$ as $A_1x = A_2x + b$
- Choose A_1 so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$
 - Requires guess, x^0

Jacobi iteration

- **Basic idea:** rewrite $Ax=b$ as $A_1x = A_2x + b$
- **Choose A_1** so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$
 - Requires guess, x^0
- **Jacobi iteration:** Choose A_1 to be diagonal matrix (main diagonal of A):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Jacobi iteration

- **Basic idea:** rewrite $Ax=b$ as $A_1x = A_2x + b$
- **Choose A_1** so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$
 - Requires guess, x^0
- **Jacobi iteration:** Choose A_1 to be diagonal matrix (main diagonal of A):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Main algorithm, easy to code!

Jacobi iteration in Fortran

- **Plan:**
 - **Set parameters:** a, b, n, tol
 - **Construct grid** x_i
 - **Construct source function, $S(x)$, initialize $T=T(x,t=0)$**
 - **Iterate using formula below**
 - **Each iteration check if $|T_k - T_{k-1}| < \text{tol}$**

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Jacobi iteration in Fortran

- **One Fortran trick: set variables to be dimension(0:N+1)**
 - $x(0)=0$, $x(N+1)=1$, $T(0)=a$, $T(N+1)=b$
 - **Then, easy to compute T_1 using:**

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Jacobi iteration in Fortran

- **One Fortran trick: set variables to be dimension(0:N+1)**
 - $x(0)=0$, $x(N+1)=1$, $T(0)=a$, $T(N+1)=b$
 - **Then, easy to compute T_1 using:**

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Core part of code (see *jacobi1s.f90*):

```
do k1=1, kmax
```

```
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
```

```
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
```

```
  T(1:n)=Tnew(1:n)      !update variable
```

```
  if (deltaT(k1)<tol) exit !check convergence criterion
```

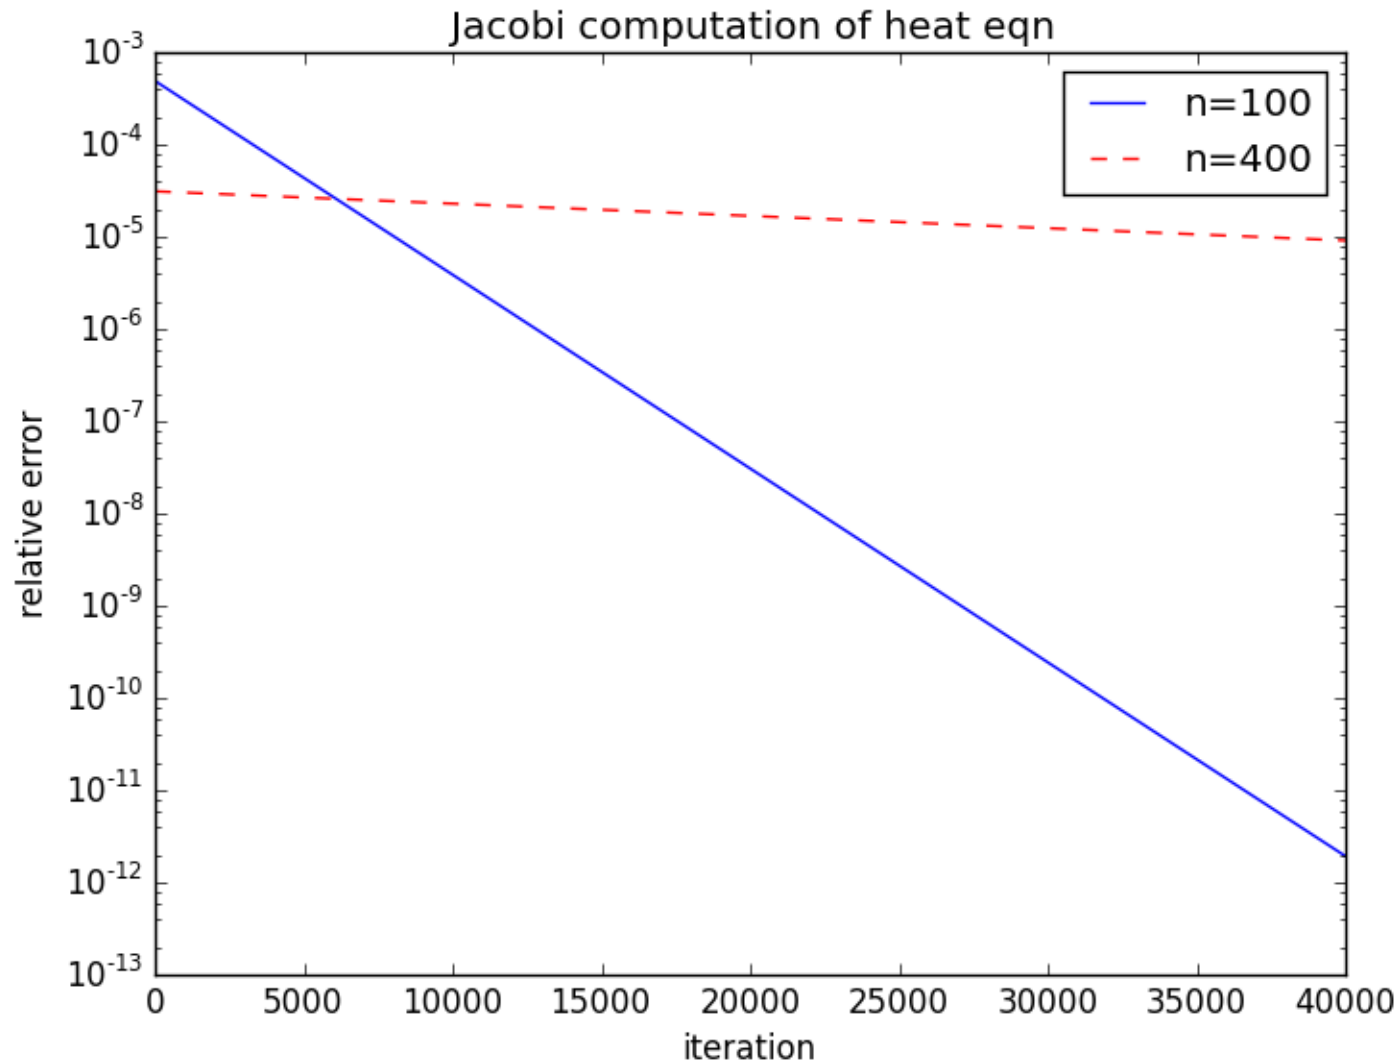
```
end do
```

Jacobi results

1. Does the solution converge?

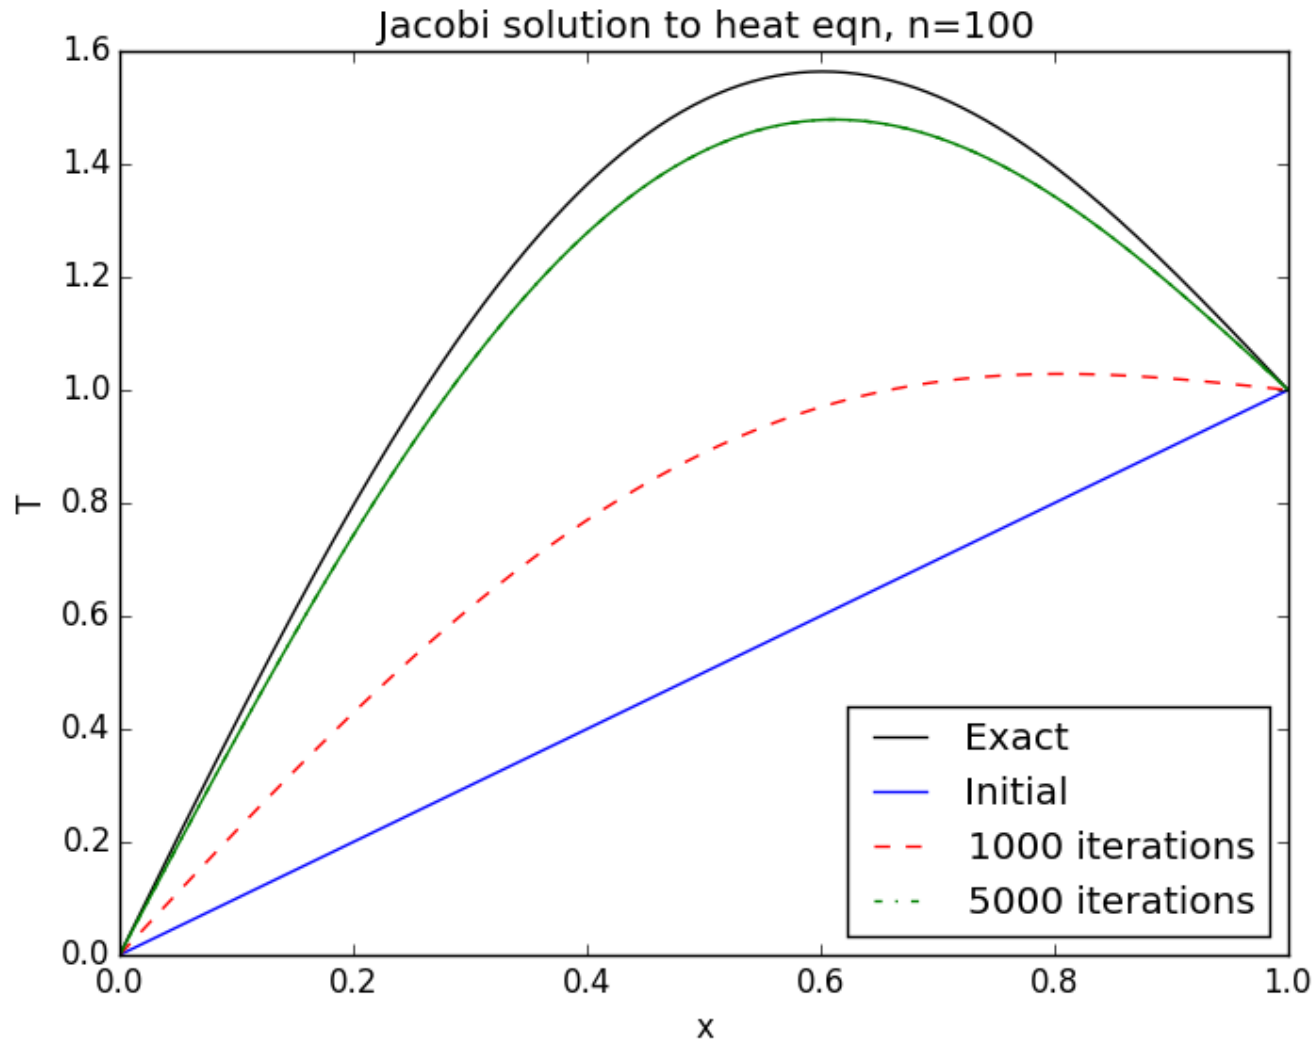
Jacobi results

1. Does the solution converge? Yes, but very slowly for large n



Jacobi results

1. Does it converge to the correct solution? Yes, can check that error $\sim \Delta x^2$




Jacobi results

- Jacobi is simplest, but *most inefficient* iterative solver
- Good illustration of basic ideas
- Better methods: Gauss-Seidel, SOR, conjugate gradient, multigrid

Algorithms and hardware


Speed-Up
Factor

Derived from Computational Methods



Speed-Up
Factor

Derived from Supercomputer Hardware



Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Serial:

```
Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
```

```
deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
```

Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
    dmax=0.d0
    !$omp parallel do reduction(max:dmax)
    do i1=1,n
        Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
        dmax = max(dmax,abs(Tnew(i1)-T(i1)))
    end do
    !$omp end parallel do
    deltaT(k1) = dmax
```

Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Serial:

```
do i1=0,n+1
    x(i1) = i1*dx
end do
!-----

!set initial condition
T = (b-a)*x + a

!set source function
S = S0*sin(pi*x)
```

Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
!$omp parallel do
do i1=0,n+1
    x(i1) = i1*dx
    T(i1) = (b-a)*x(i1) + a !set initial condition
    S(i1) = S0*sin(pi*x(i1)) !set source function
end do
!$omp end parallel do
```

Parallel Jacobi notes

- Will only see speedup with $n > \sim 20000$ (commonly seen in 2D problems)
- See *`jacobi1s_omp.f90`*, *`jacobi1_omp.py`*
- **f2py and OpenMP:** `f2py --f90flags='-fopenmp' -lgomp -c jacobi1s_omp.f90 -m j1`
- **On my laptop:**
`f2py --f90flags='-fopenmp' -L/usr/local/lib -lgomp -c jacobi1s_omp.f90 -m j1`

Time-dependent problem

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

$$0 \leq x \leq 1$$

Simple (inefficient) approach: *method of lines*

- 1. Discretize spatial variable \rightarrow $N+2$ points between 0 and 1**
- 2. Solve resulting N ODEs with solver of choice (*odeint, ode15s, ...*)**

Time-dependent problem

Again, we discretize the derivative as:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

Time-dependent problem

Again, we discretize the derivative as:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

So, we have **N ODEs**:

$$\frac{dT_i}{dt} = S_i(t) + \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}, \quad i = 1, 2, \dots, N$$

with the boundary conditions substituted in the RHS when needed.

Solving single ODE in python

- **Use *odeint* from *scipy.integrate* module to solve:**

$$\frac{dy}{dt} = -ay$$

- **Basic idea: discretize time, $t = 0, dt, \dots, N*dt$, and starting from $y(0)$ march forward in time and compute $y(dt), \dots y(N*dt)$**
- ***odeint* chooses the stepsize, dt , so that error tolerances are satisfied**
- **Need to specify:**
 - **Initial condition**
 - **Timespan for integration**
 - **A Python function which provides RHS of the ODE to *odeint***
- **Look at *ode_example.py* and lab 4**

Time-dependent problem

Solving N ODEs:

$$\frac{dT_i}{dt} = S_i(t) + \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}, \quad i = 1, 2, \dots, N$$

- Will need to provide N initial conditions when calling *odeint*.
- The python function which provides RHS to *odeint* will:
 - Take t and T_1, \dots, T_N and any other needed parameters as input
 - Return N values for dT/dt as output
- No need for Fortran for 1D problems, but may be faster for two and three dimensions.