

High Performance Computing

Autumn, 2018

Lecture 18

Notes

M3C final project: please keep an eye on *online* assignment for clarifications/corrections

Feedback: All provided marks are *provisional* and subject to rescaling by an exam committee in June

Check mpif90/mpiexec installation, both need to be in path, see notes on project

MPI datatypes

List of MPI datatypes (Fortran)

MPI Datatype

MPI_CHARACTER

MPI_COMPLEX

MPI_DOUBLE_PRECISION

MPI_INTEGER

MPI_LOGICAL

MPI_PACKED

MPI_BYTE

Fortran Datatype

CHARACTER

COMPLEX

DOUBLE PRECISION

INTEGER

LOGICAL

MPI reductions

List of MPI reductions:

<u>Reduction</u>	<u>Description</u>	<u>Datatype(s)</u>
MPI_MAX	maximum	integer, floating point
MPI_MIN	minimum	
MPI_SUM	sum	integer, floating point, complex, multilanguage types
MPI_PROD	product	
MPI LAND	logical and	logical
MPI_LOR	logical or	
MPI_LXOR	logical xor	
MPI_BAND	bitwise and	integer, byte, multilanguage types
MPI_BOR	bitwise or	
MPI_BXOR	bitwise xor	
MPI_MAXLOC	max value and location	MPI_DOUBLE_INT and such
MPI_MINLOC	min value and location	


Special datatypes needed for maxloc,minloc, see:

https://www.open-mpi.org/doc/v2.0/man3/MPI_Reduce.3.php

2-D (steady) heat equation

What is best domain decomposition?


If we have four processors, can try:



2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:




Then, parallelization is essentially same as differentiation example:

- Loop across rows
- At “boundaries”, send/recv data needed to compute second derivative
- Reduce $\max(|\Delta T|)$

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:




- But is the 1D decomposition the best?
- Want to minimize communication
- M “layers”: of a $n \times n$ grid:
 - $(M-1) \cdot n$ boundary points

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:




- But is the 1D decomposition the best?

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:




- But is the 1D decomposition the best?
- Want to minimize communication
- M “boxes”: of a $n \times n$ grid:
 - Each interior box has $2n/\sqrt{M}$ boundary points
 - Total: $2n \cdot (\sqrt{M} - 1)$ boundary points

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:



“Boundary points” are on the red dashed lines

- But is the 1D decomposition the best?
- Want to minimize communication
- M “boxes”: of a $n \times n$ grid:
 - Each box has $2n/\sqrt{M}$ boundary points
 - Total: $2n*(\sqrt{M} - 1)$ boundary points
- Boxes: less communication, but more difficult to implement!

2-D (steady) heat equation

- **MPI provides tools for creating and managing complex “topologies”**
- **For example, to create a 4 x 3 “grid” of processes:**

call `MPI_cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`

with: `ndims = 2, dims = (/4,3/), periods = (/false./,true/), reorder = .false.`

- **Here, periods sets periodic boundary conditions along the three columns**

2-D (steady) heat equation

- **MPI provides tools for creating and managing complex “topologies”**
- **For example, to create a 4 x 3 “grid” of processes:**


call `MPI_cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`

with: `ndims = 2, dims = (/4,3/), periods = (/false./true/), reorder = .false.`

- **Here, periods sets periodic boundary conditions along the three columns**
- **Other useful commands with the new communicator, new_comm:**
 - **MPI_Cart_coords: given process id, provides (i,j) coordinate**
 - **MPI_Cart_rank: given (i,j), provides id (0, 1, 2, ..., numprocs)**
- **Most useful: MPI_Cart_shift: provides id of neighboring processes in horizontal or vertical direction**
 - **Use to set up send/recv sequences needed for exchanging boundary data.**

2-D (steady) heat equation

- **Most useful:** `MPI_Cart_shift`: provides id of neighboring processes in horizontal or vertical direction
 - Use to set up send/rcv sequences needed for exchanging boundary data.
- How to decide on process grid dimensions?
- `MPI_Dims_create`:
Given number number of processes and dimensions,
outputs process grid
Dimensions (4,3 in picture →)
- e.g. 400 x 300 grid points:
4 x 3 process grid with
100 x 100 points on each grid



Synchronization (OpenMP)

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- *Barriers* keep the threads synchronized:

```
!$OMP parallel
```

```
!Some code
```

```
!$OMP barrier
```

```
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier

Synchronization (MPI)

- Same idea as in OpenMP
- Use *MPI_BARRIER* to keep processes synchronized
- *Barriers* keep the threads synchronized:

```
call MPI_BARRIER(MPI_COMM_WORLD, ierr)
```

See *gradient_p.f90* from lecture 16

Timing code

1. Lazy approach:

```
$ time mpiexec -n 2 midpointpt
```

```
real    0m0.073s
```

```
user    0m0.081s
```

```
sys 0m0.030s
```


Timing code

1. Lazy approach:

```
$ time mpiexec -n 2 midpointpt
```

```
real    0m0.073s
user    0m0.081s
sys 0m0.030s
```

2. Use *MPI_WTIME* to time particular parts of code:

```
starttime = MPI_WTIME() !***START TIMER***
!code...
!
!
endtime = MPI_WTIME() !***STOP TIMER***
print *, 'time= ',endtime - starttime, 'seconds'
```

Timing code

1. Lazy approach:

```
$ time mpiexec -n 2 midpointpt
```

```
real    0m0.073s
user    0m0.081s
sys 0m0.030s
```


2. Use *MPI_WTIME* to time particular parts of code:

```
starttime = MPI_WTIME() !***START TIMER***
!code...
!
!
endtime = MPI_WTIME() !***STOP TIMER***
print *, 'time= ',endtime - starttime, 'seconds'
```

3. But to get detailed information, use a profiler: VampirTrace, IPM, ...

Timing code

From IPM webpage (<http://ipm-hpc.sourceforge.net/>):



OpenMP vs MPI

OpenMP:

- **Advantages:**
 - Easy to implement (particularly for loops)
 - Single code for serial and parallel execution
 - Included with most compilers
- **Disadvantages:**
 - Can only be used for shared-memory computers

OpenMP vs MPI

OpenMP:

- **Advantages:**
 - Easy to implement (particularly for loops)
 - Single code for serial and parallel execution
 - Included with most compilers
- **Disadvantages:**
 - Can only be used for shared-memory computers

MPI:

- **Advantages:**
 - Can be used for distributed- or shared-memory systems
 - Distributed memory computing:
 - Each process has its own local variables
 - Programmer has detailed control over communication
- **Disadvantages:**
 - Relatively difficult to implement
 - Typically have separate serial and parallel codes

Choosing the right tools

- We have studied two programming languages and two approaches to parallelization
- When should we use Python, Fortran, Fortran+OpenMP, Fortran+MPI?

Choosing the right tools

- We have studied two programming languages and two approaches to parallelization
- When should we use Python, Fortran, Fortran+OpenMP, Fortran+MPI?
- Rules of thumb for solving PDEs (highly subjective)
 - 1D problems: interpreted language
 - 2D, single equation: compiled language
 - 2D, system of PDEs: compiled language, OpenMP (2-16 cores)
 - 3D, single equation: compiled language, OpenMP or maybe MPI (depending on problem size) (4-128 cores)
 - 3D system of PDES: compiled language + MPI (128+ cores)

Choosing the right tools

- We have studied two programming languages and two approaches to parallelization
- When should we use Python, Fortran, Fortran+OpenMP, Fortran+MPI?
- Rules of thumb for solving PDEs (highly subjective)
 - 1D problems: interpreted language
 - 2D, single equation: compiled language
 - 2D, system of PDEs: compiled language, OpenMP (2-16 cores)
 - 3D, single equation: compiled language, OpenMP or maybe MPI (depending on problem size) (4-128 cores)
 - 3D system of PDES: compiled language + MPI (128+ cores)

But remember: Matlab and Python have libraries which are essentially compiled Fortran/c codes!

Beyond M3C

- We have focused on a few problems
 - evolutionary games, neural networks, diffusion
- But ideas on parallelization can be easily generalized
- Consider extremely large datasets collected by Google (or the NSA!)
- This data *must* be processed with distributed parallel computing
- Want to find patterns, correlations, extrema
- But calculations will require data stored on different machines
 - Google developed *MapReduce* ~15 years ago for these problems

<http://research.google.com/archive/mapreduce.html>

(Google is now using something called, *Cloud Dataflow*)

Simple MapReduce example

Example: Count word occurrences

```
subroutine map(String input_key, String input_value):  
  // input_key: document name  
  // input_value: document contents  
  for each word w in input_value:  
    EmitIntermediate(w, "1");
```

Subroutine *map*: From document (*input_key*), read in data (*input_value*) count occurrences of a word

Simple MapReduce example

Example: Count word occurrences

```
subroutine map(String input_key, String input_value):  
    // input_key: document name  
    // input_value: document contents  
    for each word w in input_value:  
        EmitIntermediate(w, "1");  
  
subroutine reduce(String output_key, Iterator intermediate_values):  
    // output_key: a word  
    // output_values: a list of counts  
    int result = 0;  
    for each v in intermediate_values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```

Subroutine *map*: From document (*input_key*), read in data (*input_value*) count occurrences of a word

Subroutine *reduce*: Sum total number of occurrences computed by *map*

Simple MapReduce example

Example: Count word occurrences

```
subroutine map(String input_key, String input_value):
```

```
  // input_key: document name
```

```
  // input_value: document contents
```

```
  for each word w in input_value:
```

```
    EmitIntermediate(w, "1");
```

```
subroutine reduce(String output_key, Iterator intermediate_values):
```

```
  // output_key: a word
```

```
  // output_values: a list of counts
```

```
  int result = 0;
```

```
  for each v in intermediate_values:
```

```
    result += ParseInt(v);
```

```
  Emit(AsString(result));
```

Similar to midpoint_mpi.f90!:
map: each process assigned a subdomain, computes partial sum

reduce: total sum from partial sums

Subroutine *map*: From document (input_key), read in data (input_value) count occurrences of a word

Subroutine *reduce*: Sum total number of occurrences computed by *map*