

High Performance Computing

Autumn, 2018

Lecture 20

Notes

- Lab sessions this week are replaced with office hours
 - All office hours in MLC
- Homework 3 marks should be available Friday
- Final project marks and solutions will be available late January
- Please remember to fill out SOLE surveys (especially if you like the class!)
- I will also ask/collect (anonymous) feedback after project marks are done
- A (hopefully final) clarification on part 2.4 of final project was posted yesterday (bounds for θ^* are same as the bounds of θ)

Today

- *Apache Spark*
- *GPU computing*
- *Google TensorFlow*
- *Class overview*

Large-scale cluster computing

- Previously: discussed similarities between *mapReduce* and ideas covered in course
- If it is similar to older ideas, why is it important?
 - One important feature is *fault tolerance*
 - When running on 1000's of cores, must expect and plan for hardware failure
 - mapReduce includes features that allow a job to keep running (or easily be restarted) when individual cores fail

Large-scale cluster computing

- Previously: discussed similarities between *mapReduce* and ideas covered in course
- If it is similar to older ideas, why is it important?
 - One important feature is *fault tolerance*
 - When running on 1000's of cores, must expect and plan for hardware failure
 - mapReduce includes features that allow a job to keep running (or easily be restarted) when individual cores fail
 - Hadoop is an open-source tool which implements mapReduce on clusters (widely used for large-scale data processing)
 - Apache Spark is a newer cluster computing tool that builds on mapReduce and Hadoop

Apache Spark

- **MPI + Fortran/c** → large-scale, distributed-memory scientific computing
- **Spark + Python (or Java, Scala, Julia, R)** → large-scale distributed-memory data analysis
- ***PySpark***: use python to work with large datasets on clusters with built-in fault tolerance
- **Basic workflow:**
 - **Distribute data across cluster:** create a Resilient Distributed Dataset (RDD)
 - **Process (transform) the data in the RDD (in parallel):**
 - ***map***: apply an operation to each element of data, e.g. multiply each number by 2
 - ***filter***: extract portions of dataset that satisfy some criteria, e.g. extract all even numbers
 - **And there are other transformations as well**

Apache Spark

- **Basic workflow:**
 1. **Distribute data across cluster: create a Resilient Distributed Dataset (RDD)**
 2. **Process (transform) the data in the RDD (in parallel):**
 - *map*: apply an operation to each element of data, e.g. multiply each number by 2
 - *filter*: extract portions of dataset that satisfy some criteria, e.g. extract all even numbers
 - And there are a few other common transformations as well
 - Ultimately Original RDD is *transformed* into a new RDD
 3. **Extract key results from transformed RDD**
 - *reduce(function)*: define python function which is used to reduce data
 - e.g. sum, min, max, product (as usual)
 - but can construct any function which takes two elements on input and produces one output value (function must be commutative and associative)

Apache Spark

- **Basic workflow:**
 1. **Distribute data across cluster: create a Resilient Distributed Dataset (RDD)**
 2. **Process (transform) the data in the RDD (in parallel):**
 - *map*: apply an operation to each element of data, e.g. multiply each number by 2
 - *filter*: extract portions of dataset that satisfy some criteria, e.g. extract all even numbers
 - And there are a few other common transformations as well
 - Ultimately Original RDD is *transformed* into a new RDD
 3. **Extract key results from transformed RDD**
 - *reduce(function)*: define python function which is used to reduce data
 - *take*: collect 1st n elements
 - *collect*: returns all n elements (similar to gather)

Apache Spark

Example

- After downloading pre-built package for Spark:
<http://spark.apache.org/downloads.html>)

- Launch pyspark on 2 local threads:

```
$ ./bin/pyspark --master local[2]
```

- Create data to be analyzed:

```
>>> L = range(20)
```

- Create RDD:

```
>>> D = sc.parallelize(L)
```

```
>>> D
```

```
ParallelCollectionRDD[0] at parallelize at PythonRDD.scala:423
```

- D is automatically distributed across the 2 threads

Apache Spark

Example

- **Create data to be analyzed:**

```
>>> L = range(20)
```

- **Create RDD:**

```
>>> D = sc.parallelize(L)
```

```
>>> D
```

```
ParallelCollectionRDD[0] at parallelize at PythonRDD.scala:423
```

- **Define filter function (checks if input is even):**

```
>>> def is_even(x):  
    return x % 2 == 0
```

- **Filter data:**

```
>>> Dnew = D.filter(is_even)
```

- **Collect data:**

```
>>> Dnew.collect()
```

```
[0, 2, 4, 6, 8, 10, 12, 14, 16, 18]
```

Apache Spark

Example

- Can also reduce data using:
 - `Dnew_reduce(function)` where function has been specified
- Or by using a intrinsic reduction:

```
>>> Dnew.count()
10
>>> Dnew.sum()
90
```

All of this is done in parallel!

Apache Spark

Notes: Test your computation locally, but run on cluster:

- There are tools specifically designed for building “spark clusters”:
<http://spark.apache.org/docs/latest/cluster-overview.html>
- Can use scripts to launch a EC2 spark cluster
- Libraries for spark cluster computing:
 - GraphX: graph processing
 - Mllib: machine learning
<http://spark.apache.org/docs/latest/mllib-guide.html>
 - Lapack is included in *Breeze*
 - Spark SQL, Spark Streaming

- **Machine learning snippet:**

```
from pyspark.mllib.classification import SVMWithSGD, SVMModel
from pyspark.mllib.regression import LabeledPoint
model = SVMWithSGD.train(parsedData, iterations=100)
```

Aside: Lambda functions in Python

Note: often convenient to use “lambda functions”

```
>>> is_even = lambda x: x % 2 == 0
```

- **Command above is equivalent to function on earlier slide**
- **Similarly these “regular” and lambda functions are equivalent:**

```
>>> def prod(x,y):  
    return x*y
```

```
>>> prod = lambda x,y: x*y
```

GPU computing

- Graphics cards are highly efficient, very powerful
 - 100s or even 1000s of compute cores
 - Energy efficient – see the green500 list: <https://www.top500.org/green500>
 - Nvidia is the leader in GPU-accelerated computing
 - Why are GPUs so advanced?

GPU computing

- Graphics cards are highly efficient, very powerful
 - 100s or even 1000s of compute cores
 - Energy efficient – see the green500 list: <https://www.top500.org/green500>
 - Nvidia is the leader in GPU-accelerated computing
 - Why are GPUs so advanced?
 - Video games are a billion dollar industry
 - Realistic graphics require enormous computational power

GPU computing

- This computational power has been adapted for scientific computing
- Very useful for small-to-medium sized jobs with scope for parallelization
- Main limitation is memory

Nvidia Tesla P100 specs:

SPECIFICATIONS

GPU Architecture	NVIDIA Pascal
NVIDIA CUDA® Cores	3584
Double-Precision Performance	5.3 TeraFLOPS
Single-Precision Performance	10.6 TeraFLOPS
Half-Precision Performance	21.2 TeraFLOPS
GPU Memory	16 GB CoWoS HBM2
Memory Bandwidth	732 GB/s
Interconnect	NVIDIA NVLink
Max Power Consumption	300 W
ECC	Native support with no capacity or performance overhead
Thermal Solution	Passive
Form Factor	SXM2
Compute APIs	NVIDIA CUDA, DirectCompute, OpenCL™, OpenACC



TeraFLOPS measurements with NVIDIA GPU Boost™ technology

GPU computing

- **Three main approaches:**
 - **Libraries (cuBLAS, cuFFT, ...)**
 - **Directives (OpenACC), similar to OpenMP**
 - **Programming (CUDA Python, openCL, ...)**
- **Programming is not easy – have to explicitly move data between CPU and GPU and be precise with variable types and sizes**
- **Has become extremely popular for machine learning, neural networks (e.g. TensorFlow)**

TensorFlow

- TensorFlow is a general library for efficient array (“tensor”) computations on directed graphs.
- Many libraries built on this very general framework (e.g. Keras for deep learning)
- Can use a Python interface
- Some standard examples to try out (with neural networks)
 - Character recognition
 - Image recognition



Main class takeaways

- Solving large, complex problems:
 - Break problem into smaller parts:
 - Python modules with multiple functions
 - Fortran modules
 - Distribute tasks across threads/processes

Main class takeaways

- Solving large, complex problems:
 - Break problem into smaller parts:
 - Python modules with multiple functions
 - Fortran modules
 - Distribute tasks across threads/processes
 - Manage smaller parts with appropriate tools
 - git, makefiles
 - OpenMP, MPI

Main class takeaways

- Solving large, complex problems:
 - Break problem into smaller parts:
 - Python modules with multiple functions
 - Fortran modules
 - Distribute tasks across threads/processes
 - Manage smaller parts with appropriate tools
 - git, makefiles
 - OpenMP, MPI
 - Carefully test each part
 - Package test routines with the code
 - Unit testing tools (e.g. nose) can automatically run through test routines when code is changed

Main class takeaways

- **Choose the right tool for the right problem:**
 - **Interpreted vs. compiled languages**
 - **General purpose vs. scientific**
 - **Serial vs. parallel**
 - **Shared memory vs. distributed memory**
 - **Libraries vs. writing your own code**

Libraries

- **Avoid writing code whenever possible! Don't re-invent the wheel.**
- **Many powerful libraries available: lapack, boost (C++), fftw, NAG, ...**
 - **To use: 1. call subroutine/function in code**
2. link to library when compiling `gfortran -o a.exe a.f90 -l lapack`
- **Also have *parallel* libraries: scalapack, petsc**
- **These are not easy to use, but much better than writing your own code.**
- **E.g. scalapack has routines for parallel:**
 - **linear systems of equations**
 - **linear least squares**
 - **standard eigenvalue problems**
 - **singular value decomposition**

Main class takeaways

- It is important to know useful programming languages (this is what HR looks for)
- It is also important to know how to learn new programming languages (this is what experts look for)
 - What is the basic structure of (many) programming languages?

Main class takeaways

- It is important to know useful programming languages (this is what HR looks for)
- It is also important to know how to learn new programming languages (this is what experts look for)
 - What is the basic structure of (many) programming languages?
 - 10-lecture c course taught at Imperial:

The course covers:

- Different number types in C (Integers and Floating Point).
- Operators, operands and their precedence.
- Conversions and casts.
- Mathematical expressions.
- Statements: choice, while, do-while, switch, for loops...
- Functions.
- Pointers, arrays and matrices.
- Characters, strings and interacting with the console.
- Reading and writing files.
- Optimisation and Debugging.
- Scientific C-Libraries and their uses: NAG, GSL, etc.
- An Introduction to Parallel Computing.
- C++ and other languages.

Main class takeaways

- **Lectures: Computing**
- **Projects: Computing + Science + Math**
 - **Next Spring: Computational PDEs, Scientific Computing**